

Principles of Database Systems (CS307)

Lecture 6: Advanced SQL

Zhong-Qiu Wang

Department of Computer Science and Engineering
Southern University of Science and Technology

- Most contents are from slides made by Stéphane Faroult and the authors of Database System Concepts (7th Edition).
- Their original slides have been modified to adapt to the schedule of CS307 at SUSTech.
- The slides are largely based on the slides provided by Dr. Yuxin Ma

Announcements

- Second assignment is out, due date: **21st Oct., Tuesday, 10pm**
 - Do not miss the deadline, or you will receive reduced scores

Function

Built-in Functions

- Most DBMS provides a series of built-in functions
 - E.g., Scalar function, aggregation function, window function



```
round(3.141592, 3) -- 3.142  
trunc(3.141592, 3) -- 3.141
```



```
upper('Citizen Kane')  
lower('Citizen Kane')  
substr('Citizen Kane', 5, 3) -- 'zen'  
trim(' Ooops ') -- 'Ooops'  
replace('Sheep', 'ee', 'i') -- 'Ship'
```

`count(*)/count(col), min(col), max(col), stddev(col), avg(col)`

`<function> over (partition by <col_p> order by <col_o1, col_o2, ...>)`

- `<function>`: we can apply (1) ranking window functions, or (2) aggregation functions
- `partition by`: specify the column for grouping
- `order by`: specify the column(s) for ordering in each group

Self-defined Function

- Sometimes the built-in functions cannot fulfill our requirements
 - And the power of declarative language (SQL) is not strong enough
- Most DBMS implement a **built-in, SQL-based programming language**
 - A **procedural extension** to SQL

Procedural vs. Declarative

- Two different programming paradigms
 - Imperative programming (命令式编程)
 - Describe the algorithms step-by-step (i.e., how to do)
 - Procedural (过程式) : C (and many other legacy languages)
 - Object-oriented: Java
 - Declarative programming (声明式编程)
 - Describe the result without specifying the detailed steps (i.e., what to do)
 - (Pure) declarative: SQL, Regular Expressions, Markup (HTML, XML), CSS
 - Functional: Scheme, Haskell, Scala, Erlang
 - Logic programming: Prolog

Procedural vs. Declarative

- E.g., How can we get a cup of tea?

- In a procedural way:

1. Get a cup
2. Get some tea
3. Get some hot water
4. Put tea into the cup
5. Pour hot water into the cup
6. return tea;



- In a declarative way:

<a cup of tea/>

- You don't really need to know how to make a cup of tea
 - The system can do it in a black-box manner



大佬喝茶

Procedural vs. Declarative

- E.g., Find all Chinese movies before 1990 in the movies table?

- In a procedural way:

1. Read the movies table into the memory
2. For each row *i* in the table, repeat:
 - 2.1 In row *i*, read the value of the column “country”
 - 2.2 if ...

-
- In a declarative way: `select * from movies where country = 'cn' and year_released < 1990`
 - You don't really need to know how to filter the table
 - The DBMS system can do it in a black-box manner

Procedural vs. Declarative

- Benefits in declarative languages
 - No need to understand the details
 - The systems take in charge of all the details
 - Easier to use than imperative programming
 - More user-friendly
- Problems in declarative languages
 - Cannot specify the control flow of a program
 - If there is no such command as <a cup of tea/>, you need to create it by yourself

Procedural Extension to SQL

- Many DBMS products provide a **proprietary** procedural extension to the standard SQL
 - Transact-SQL (T-SQL) 
 - PL/SQL 
 - PL/PGSQL 
 - (No specific name) 
 - (Not supported) 

Function in (Postgre)SQL

- Example: Display the full name for people with “von”
 - When introducing **update**, we have modified the names starting with “von” into “... (von)” for ordering
 - **von Neumann** -> **Neumann (von)**

| | peopleid | first_name | surname | born | died | gender |
|----|----------|------------|----------------------|------|--------|--------|
| 1 | 16439 | Axel | Ambesser (von) | 1910 | 1988 | M |
| 2 | 16440 | Daniel | Bargen (von) | 1950 | 2015 | M |
| 3 | 16441 | Eduard | Borsody (von) | 1898 | 1970 | M |
| 4 | 16442 | Suzanne | Borsody (von) | 1957 | <null> | F |
| 5 | 16443 | Tomas | Brömssen (von) | 1943 | <null> | M |
| 6 | 16444 | Erik | Detten (von) | 1982 | <null> | M |
| 7 | 16445 | Theodore | Eltz (von) | 1893 | 1964 | M |
| 8 | 16446 | Gunther | Fritsch (von) | 1906 | 1988 | M |
| 9 | 16447 | Katja | Garnier (von) | 1966 | <null> | F |
| 10 | 16448 | Harry | Meter (von) | 1871 | 1956 | M |
| 11 | 16449 | Jenna | Oÿ (von) | 1977 | <null> | F |
| 12 | 16450 | Alicia | Rittberg (von) | 1993 | <null> | F |
| 13 | 16451 | Daisy | Scherler Mayer (von) | 1966 | <null> | F |
| 14 | 16452 | Gustav | Seyffertitz (von) | 1862 | 1943 | M |

Function in (Postgre)SQL

- If we simply concatenate the first name and the last name, it looks like this:
 - A little bit weird format (a trailing “von”)



```
select first_name || ' ' || surname  
from people  
where surname like '%(von)' ;
```

| | ?column? |
|----|----------------------------|
| 1 | Axel Ambesser (von) |
| 2 | Daniel Bargen (von) |
| 3 | Eduard Borsody (von) |
| 4 | Suzanne Borsody (von) |
| 5 | Tomas Brömssen (von) |
| 6 | Erik Detten (von) |
| 7 | Theodore Eltz (von) |
| 8 | Gunther Fritsch (von) |
| 9 | Katja Garnier (von) |
| 10 | Harry Meter (von) |
| 11 | Jenna Oÿ (von) |
| 12 | Alicia Rittberg (von) |
| 13 | Daisy Scherler Mayer (von) |
| 14 | Gustav Seyffertitz (von) |

Function in (Postgre)SQL

- Question: How can we restore the format into “`first_name von surname`”?
 - String operations
 - i.e., `Neumann (von)` -> `von Neumann`

Function in (Postgre)SQL

- Question: How can we restore the format into “first_name von surname”?
 - String operations
 - i.e., Neumann (von) -> von Neumann

```
select case
    when first_name is null then ''
    else first_name || ' '
end || case position('(' in surname)
    when 0 then surname
    else trim(')' from substr(surname, position('(' in surname) + 1))
        ||
        || trim(substr(surname, 1, position('(' in surname) - 1))
end
from people
where surname      like '%(von)';
```

Function in (Postgre)SQL

- Question: How can we restore the format into “first_name von surname”?
 - String operations




Then, how can we store this part to reuse it in the future?

```
case
when first_name is null then ''
else first_name || ' '
end || case position('(' in surname)
when 0 then surname
else trim(')' from substr(surname, position('(' in surname) + 1))
|| ' '
|| trim(substr(surname, 1, position('(' in surname) - 1))
end
from people
where surname ... like '%(von)';
```

Function in (Postgre)SQL

- Copy and paste is not a good habit
 - Whenever you have painfully written something as complicated, which is pretty generic, you'd rather not copy and paste the code every time you need it




Function in (Postgre)SQL

- Stored for reuse
 - In PostgreSQL, we can store the expression and reuse it in another context
- Self-defined Function
 - `create function`



```
CREATE [OR REPLACE] FUNCTION function_name (arguments)
RETURNS return_datatype AS $variable_name$
DECLARE
    declaration;
    [...]
BEGIN
    < function_body >
    [...]
    RETURN { variable_name | value }
END; LANGUAGE plpgsql;
```



```
CREATE [ OR REPLACE ] FUNCTION
    name ( [ [ argmode ] [ argname ] argtype [ { DEFAULT | = } default_expr ] [, ...] ]
    [ RETURNS rettype
    | RETURNS TABLE ( column_name column_type [, ...] ) ]
    { LANGUAGE lang_name
    | TRANSFORM { FOR TYPE type_name } [, ... ]
    | WINDOW
    | { IMMUTABLE | STABLE | VOLATILE }
    | [ NOT ] LEAKPROOF
    | { CALLED ON NULL INPUT | RETURNS NULL ON NULL INPUT | STRICT }
    | { [ EXTERNAL ] SECURITY INVOKER | [ EXTERNAL ] SECURITY DEFINER }
    | PARALLEL { UNSAFE | RESTRICTED | SAFE }
    | COST execution_cost
    | ROWS result_rows
    | SUPPORT support_function
    | SET configuration_parameter { TO value | = value | FROM CURRENT }
    | AS 'definition'
    | AS 'obj_file', 'link_symbol'
    | sql_body
} ...
```

...or, a simpler version

Function in (Postgre)SQL

- How do we rewrite the name conversion expression into a function?

```
create function full_name(p_fname varchar, p_sname varchar)
returns varchar
as $$ 
begin
    return case
        when p_fname is null then ''
        else p_fname || ' '
    end || case position('(' in p_sname)
        when 0 then p_sname
        else trim(')' from substr(p_sname, position('(' in p_sname) + 1))
            ||
            || trim(substr(p_sname, 1, position('(' in p_sname) - 1))
    end;
end;
$$ language plpgsql;
```

Function in (Postgre)SQL

- How do we rewrite the name conversion expression into a function?



Function name and the parameter list

- Format for variables and parameters: [name] [type]

```
create function full_name(p_fname varchar, p_sname varchar)
returns varchar
as $$
begin
    return case
        when p_fname is null then ''
        else p_fname || ' '
    end || case position('('' in p_sname)
        when 0 then p_sname
        else trim(')' from substr(p_sname, position('('' in p_sname) + 1))
            ||
            || trim(substr(p_sname, 1, position('('' in p_sname) - 1))
    end;
end;
$$ language plpgsql;
```

Function in (Postgre)SQL

- How do we rewrite the name conversion expression into a function?

```
create function full_name(p_fname varchar, p_sname varchar)
    returns varchar Return type
    as $$
begin
    return case
        when p_fname is null then ''
        else p_fname || ' '
    end || case position('('' in p_sname)
        when 0 then p_sname
        else trim(')' from substr(p_sname, position('('' in p_sname) + 1))
            ||
            || trim(substr(p_sname, 1, position('('' in p_sname) - 1))
    end;
end;
$$ language plpgsql;
```

Function in (Postgre)SQL

- How do we rewrite the name conversion expression into a function?

Body {

```
create function full_name(p_fname varchar, p_sname varchar)
returns varchar
as $$
begin
    return case
        when p_fname is null then ''
        else p_fname || ' '
    end || case position('('' in p_sname)
        when 0 then p_sname
        else trim(')' from substr(p_sname, position('(' in p_sname) + 1))
        || ''
        || trim(substr(p_sname, 1, position('(' in p_sname) - 1))
    end;
end;
$$ language plpgsql;
```

Function in (Postgre)SQL

- How do we rewrite the name conversion expression into a function?

```
create function full_name(p_fname varchar, p_sname varchar)
returns varchar
as $$

begin
    return case
        when p_fname is null then ''
        else p_fname || ' '
    end || case position('(' in p_sname)
        when 0 then p_sname
        else trim(')' from substr(p_sname, position('(' in p_sname) + 1))
        || ' '
        || trim(substr(p_sname, 1, position('(' in p_sname) - 1))
    end;
end;
$$ language plpgsql;
```

A very simple body: return the value of an expression

Function in (Postgre)SQL

- How do we rewrite the name conversion expression into a function?

```
create function full_name(p_fname varchar, p_sname varchar)
returns varchar
as $$begin
    return case
        when p_fname is null then ''
        else p_fname || ' '
    end || case position('('' in p_sname)
        when 0 then p_sname
        else trim(')' from substr(p_sname, position
            || ''
            || trim(substr(p_sname, 1, position('('' in
                end;
end;
$$ language plpgsql;
```

A very simple body: return the value of an expression

Procedural extensions provide all the features in a true (procedural) programming languages, such as:

- Variables
- Conditions
- Loops
- Arrays
- Error management
- ...

Function in (Postgre)SQL

- How do we rewrite the name conversion expression into a function?

```
create function full_name(p_fname varchar, p_sname varchar)
returns varchar
as $$ 
begin
    return case
        when p_fname is null then ''
        else p_fname || ' '
    end || case position('('' in p_sname)
        when 0 then p_sname
        else trim(')' from substr(p_sname, position('('' in p_sname) + 1))
        || ''
        || trim(substr(p_sname, 1, position('('' in p_sname) - 1))
    end;
end;
$$ language plpgsql;
```

Language Type

PostgreSQL supports 4 procedural languages: PL/pgSQL, PL/Tcl, PL/Perl, and PL/Python

- Tcl, Perl, and Python are famous scripting languages in case you don't know

Function in (Postgre)SQL

- How do we rewrite the name conversion expression into a function?

```
create function full_name(p_
returns varchar
as $$
begin
    return case
        when p_fname is null
        else p_fname || ' '
    end || case position('('' in p_sname)
        when 0 then p_sname
        else trim(')' from substr(p_sname, position('('' in p_sname) + 1))
        || ''
        || trim(substr(p_sname, 1, position('('' in p_sname) - 1))
    end;
end;
$$ language plpgsql
```

```
create function append_test(p_code varchar)
returns varchar
as $$
    if p_code == 'cn':
        return 'China'
    else:
        return 'not China'
$$ language plpython3u;
```

Yes, we can even use Python to write functions



Language Type

PostgreSQL supports 4 procedural languages:

PL/pgSQL, PL/Tcl, PL/Perl, and PL/Python

- Tcl, Perl, and Python are famous scripting languages in case you don't know

Function in (Postgre)SQL

- Once your function is created, you can use it as if it were any built-in function.



```
select full_name(first_name, surname)
from people
where surname like '%(von)';
```

Function in (Postgre)SQL

- We can run `select` queries in functions
 - Example: design a function “`get_country_name`” to transform the country codes into country names based on the `countries` table

```
create function get_country_name(p_code varchar)
returns countries.country_name%type
as $$          i.e., same type as countries.country_name
declare          v_name countries.country_name%type;
begin
    select country_name
    into v_name
    from countries
    where country_code = p_code;
    return v_name;
end;
$$ language plpgsql;
```

```
select get_country_name(country) from movies;
```

Function in (Postgre)SQL

- We can run `select` queries in functions
 - Example: design a function “`get_country_name`” to transform the country codes into country names based on the `countries` table

```
create function get_country_name(p_code varchar)
returns countries.country_name%type
as $$
declare
    v_name countries.country_name%type;
begin
    select country_name
    into v_name
    from countries
    where country_code = p_code;
    return v_name;
end;
$$ language plpgsql;
```

```
select get_country_name(country) from movies;
```

... seems to be an easy way to get rid of join operations?

```
select c.country_name
from countries c join movies m
on c.country_code = m.country;
```

Function in (Postgre)SQL

- A “look-up function” forces a “one row at a time” join which in most cases will be costly



```
select get_country_name(country) from movies;
```

For each row in movies, the select query in `get_country_name()` is executed once

More to Read

- We may not cover all the details in functions in the theoretical session, so here are some more materials on procedural programming in PostgreSQL:
 - Lab tutorial on Functions
 - Please read it before your next lab sessions
 - Chapter 5.2 “Functions and Procedures,” Database System Concepts (7th Edition)
 - Chapter 43 “PL/pgSQL,” PostgreSQL Documentation
 - <https://www.postgresql.org/docs/current/plpgsql.html>

Procedures

Functions vs. Procedures

- Generally,
 - “Function” comes from mathematics
 - ... which calculates a value with a given input (or to say, map a value to another)
 - Thus, functions always have a return value
 - “Procedure” comes from programming
 - ... which is used to describe a set of instructions that will be executed in order
 - ... and does NOT (necessarily) have a return value
- However,
 - Sometimes, the two terms are interchangeably (used for representing the same thing)
 - e.g., procedures are called functions as well
 - Be careful when seeing both terms
 - Always identify the exact meaning of each term and see whether they have different or the same meaning(s)

Functions and Procedures in (Postgre)SQL

- It follows the general definition of functions and procedures
 - **Function**: return a value
 - **Procedure**: return **NO** value
- However,
 - For some historical reasons, PostgreSQL actually has no implementation specifically for procedures
 - It shares the same mechanism with functions
 - Treats procedures as **void functions**
 - * But for some other database systems, there are separate implementations for functions and procedures

When to Use Procedures

- For business logics
 - One requirement may need a series of SQL querys and statements
 - Transactions may be used
 - Example: Insert a new movie into the databases
 - movies table
 - Basic information for the movie
 - countries table
 - Transformation between country names and codes
 - people table
 - new actors / directors
 - credits table
 - new credit information
 - Problem: Update all the tables? Input validation? Code reuse? Security?

When to Use Procedures

- To add a movie:
 - We may have a series queries to execute when inserting only one movie
 - How about one call for all the processes?
 - Benefit 1: Network overhead
 - When running multiple queries, you are going to waste time chatting over the network with the remote server
 - Benefit 2: Security
 - Prevent users from modifying data otherwise than by calling carefully written and well tested procedures
 - Ensure that users can only modify data via carefully written and well tested procedures

Example: Adding a New Movie

- The information provided for a new movie:
 - Title
 - Year
 - Country Name
 - Note: Name, not code. Country codes are not user-friendly
 - E.g. Which country does “at” represent? “al”? “ma”? “li”?
 - Director
 - Actor 1
 - Actor 2
 - Let’s assume that only one director and at most two actors are allowed
 - It will be more difficult when the number of people are flexible

A Typical Process

- Insert and check the values, constraints, existence, duplicates, etc
 - A series of inter-related statements



```
select country_code from countries  
... -- Look up the country code
```

```
insert into movies  
... -- Insert a row in the movies table
```

```
select peopleid from people  
...  
insert into credits  
... -- Director
```

```
select peopleid from people  
...  
insert into credits  
... -- Actor 1
```

```
select peopleid from people  
...  
insert into credits  
... -- Actor 2
```

A Typical Process

- How can we pack them into a single execution unit?
 - Minimize communication between client program and database server
 - Client program = DataGrip, psql, ...



```
select country_code from countries  
... -- Look up the country code
```

```
insert into movies  
... -- Insert a row in the movies table
```

```
select peopleid from people  
...  
insert into credits  
... -- Director
```

```
select peopleid from people  
...  
insert into credits  
... -- Actor 1
```

```
select peopleid from people  
...  
insert into credits  
... -- Actor 2
```

insert into select

- One thing to optimize: insert into ... select



-- when the arities of table 1 and 2 are the same:

```
insert into table2
select * from table1
where condition;
```

-- only insert specific columns

```
insert into table2 (column1, column2, column3, ...)
select column1, column2, column3, ...
from table1
where condition;
```



Optimize the Insertion

- Use `insert into select`

```
insert into movies ...
select country_code, ...
from countries
...

-- insert the director
-- by looking up the people table
insert into credits ...
select peopleid, 'D', ...
from people
...

-- insert the first actor
-- by looking up the people table
insert into credits ...
select peopleid, 'A', ...
from people
...

-- insert the second actor
-- by looking up the people table
insert into credits ...
select peopleid, 'A', ...
from people
...
```

Further Optimize the Insertion

- Use `insert into select`
- Combine the queries of people

● ● ●
`insert into movies ...
select country_code, ...
from countries`

...

*-- insert the director
-- by looking up the people table*
`insert into credits ...
select peopleid, 'D', ...
from people`

...

*-- insert the first actor
-- by looking up the people table*
`insert into credits ...
select peopleid, 'A', ...
from people`

...

*-- insert the second actor
-- by looking up the people table*
`insert into credits ...
select peopleid, 'A', ...
from people`

...

● ● ●
`insert into movies ...
select country_code, ...
from countries`

...


*-- insert all three people
-- together*
`insert into credits ...
select peopleid, ...
from (select director, 'D', ...
union all
select actor1, 'A', ...
union all
select actor2, 'A', ...) a
inner join people`

...



Further Optimize the Insertion

- Use `insert into select`
- Combine the queries of people



The Procedure



```
create function movie_registration
    (p_title      varchar,
     p_country_name varchar,
     p_year       int,
     p_director_fn  varchar,
     p_director_sn  varchar,
     p_actor1_fn   varchar,
     p_actor1_sn   varchar,
     p_actor2_fn   varchar,
     p_actor2_sn   varchar)
returns void
as $$

declare
    n_rowcount int;
    n_movieid int;
    n_people int;
begin
    insert into movies(title, country, year_released)
        select p_title, country_code, p_year
        from countries
        where country_name = p_country_name;
    get diagnostics n_rowcount = row_count;

    if n_rowcount = 0
    then
        raise exception 'country not found in table COUNTRIES';
    end if;
```

```
n_movieid := lastval();
select count(surname)
into n_people
from (select p_director_sn as surname
      union all
      select p_actor1_sn as surname
      union all
      select p_actor2_sn as surname) specified_people
where surname is not null;

insert into credits(movieid, peopleid, credited_as)
select n_movieid, people.peopleid, provided.credited_as
from (select coalesce(p_director_fn, '*') as first_name,
            p_director_sn as surname,
            'D' as credited_as
         union all
         select coalesce(p_actor1_fn, '*') as first_name,
            p_actor1_sn as surname,
            'A' as credited_as
         union all
         select coalesce(p_actor2_fn, '*') as first_name,
            p_actor2_sn as surname,
            'A' as credited_as) provided
inner join people
on people.surname = provided.surname
and coalesce(people.first_name, '*') = provided.first_name
where provided.surname is not null;

get diagnostics n_rowcount = row_count;
if n_rowcount != n_people
then
    raise exception 'Some people couldn''t be found';
end if;
end;
$$ language plpgsql;
```

The Procedure



```
create function movie_registration
    (p_title      varchar,
     p_country_name varchar,
     p_year       int,
     p_director_fn  varchar,
     p_director_sn  varchar,
     p_actor1_fn   varchar,
     p_actor1_sn   varchar,
     p_actor2_fn   varchar,
     p_actor2_sn   varchar)
returns void
as $$
declare
    n_rowcount int;
    n_movieid int;
    n_people int;
begin
    insert into movies(title, country, year_released)
        select p_title, country_code, p_year
        from countries
        where country_name = p_country_name;
    get diagnostics n_rowcount = row_count;

    if n_rowcount = 0
    then
        raise exception 'country not found in table COUNTRIES';
    end if;
```

Check whether a row in
movies is correctly inserted

```
n_movieid := lastval();
select count(surname)
into n_people
from (select p_director_sn as surname
      union all
      select p_actor1_sn as surname
      union all
      select p_actor2_sn as surname) specified_people
where surname is not null;

insert into credits(movieid, peopleid, credited_as)
select n_movieid, people.peopleid, provided.credited_as
from (select coalesce(p_director_fn, '*') as first_name,
            p_director_sn as surname,
            'D' as credited_as
         union all
         select coalesce(p_actor1_fn, '*') as first_name,
            p_actor1_sn as surname,
            'A' as credited_as
         union all
         select coalesce(p_actor2_fn, '*') as first_name,
            p_actor2_sn as surname,
            'A' as credited_as) provided
inner join people
on people.surname = provided.surname
and coalesce(people.first_name, '*') = provided.first_name
where provided.surname is not null;

get diagnostics n_rowcount = row_count;
if n_rowcount != n_people
then
    raise exception 'Some people couldn''t be found';
end if;
end;
$$ language plpgsql;
```

Check whether rows in
credits are correctly inserted

Calling Procedures

- In PostgreSQL
 - We can call the procedure interactively by calling it from a SELECT statement (that will return nothing)
 - We can also call a procedure from another procedure



```
select movie_registration('The Adventures of Robin Hood',
                          'United States', 1938,
                          'Michael', 'Curtiz',
                          'Errol', 'Flynn',
                          null, null);
```



```
perform movie_registration('The Adventures of Robin Hood',
                           'United States', 1938,
                           'Michael', 'Curtiz',
                           'Errol', 'Flynn',
                           null, null);
```