

Principles of Database Systems (CS307)

Lecture 14 part 1: Query Processing

Zhong-Qiu Wang

Department of Computer Science and Engineering
Southern University of Science and Technology

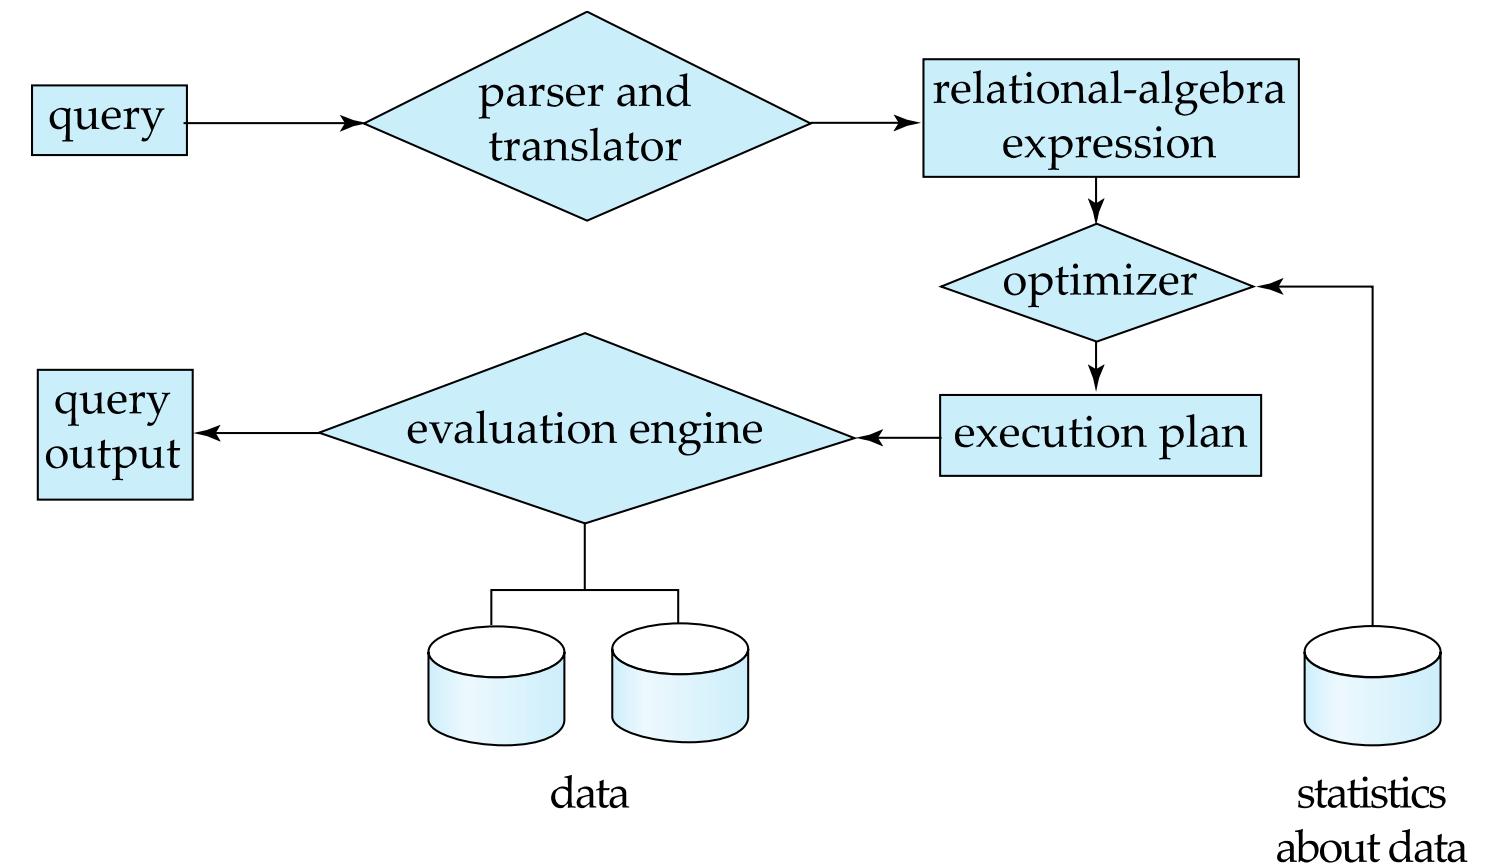
- Most contents are from slides made by Stéphane Faroult and the authors of Database System Concepts (7th Edition).
- Their original slides have been modified to adapt to the schedule of CS307 at SUSTech.
- The slides are largely based on the slides provided by Dr. Yuxin Ma

Announcements

- Final exam will be on Jan. 7, 2026 (Wed), 14:00-16:00, 三教106/107

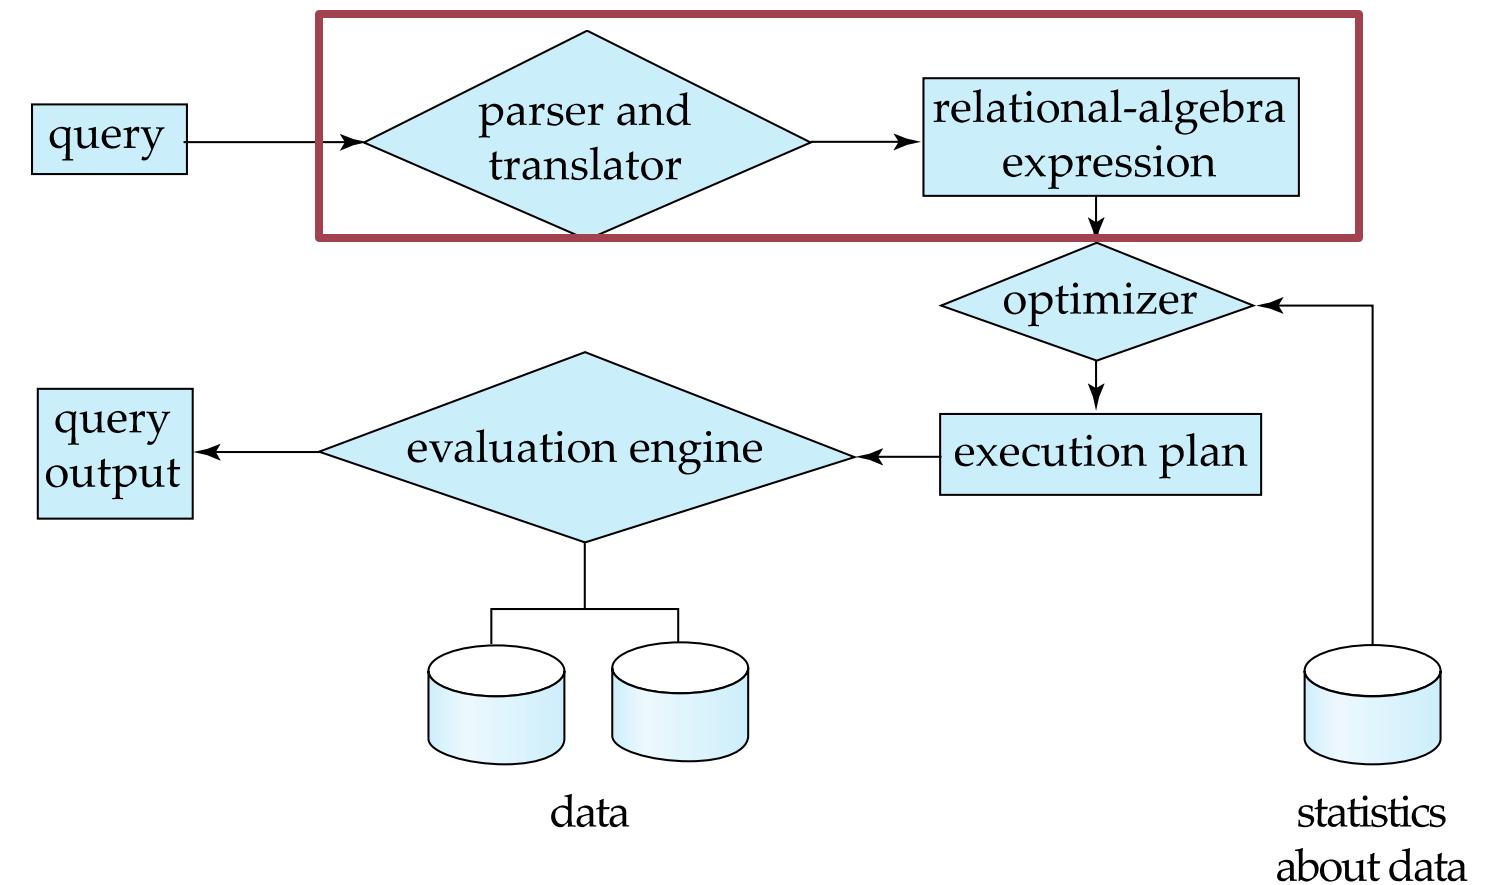
Basic Steps in Query Processing

- Parsing and Translation
- Optimization
- Evaluation



Basic Steps in Query Processing

- Parsing and Translation
 - Translate the query into its internal form
 - The internal form is then translated into **relational algebra**
 - Parser **checks syntax and verifies relations**



Basic Steps in Query Processing

- Optimization

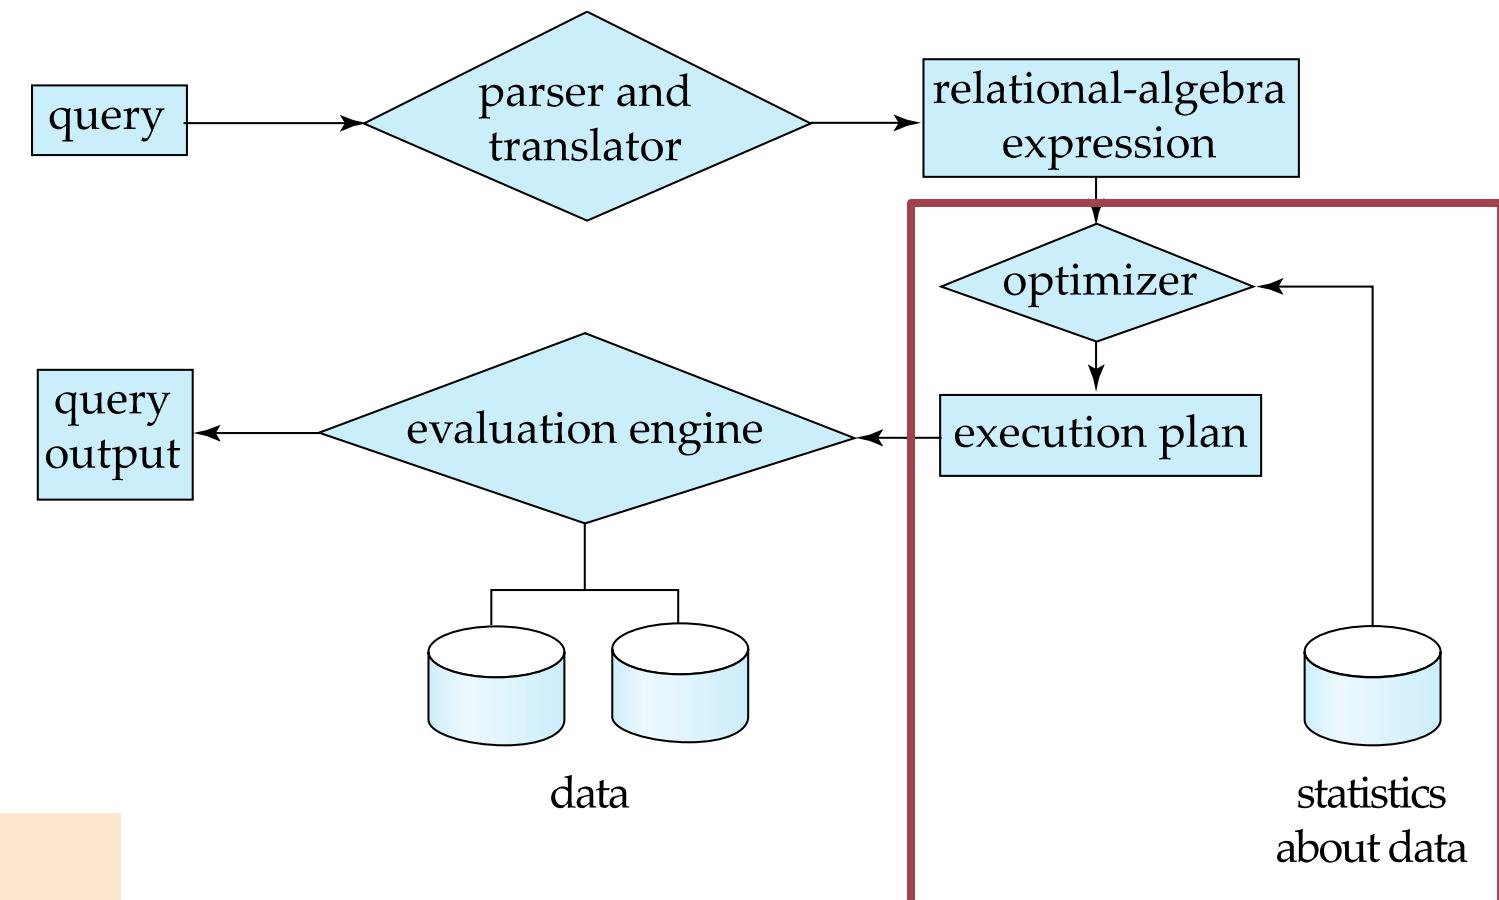
- A relational algebra expression may have many equivalent expressions
- E.g.,

$$\sigma_{\text{salary} < 75000}(\Pi_{\text{salary}}(\text{instructor}))$$

is equivalent to

$$\Pi_{\text{salary}}(\sigma_{\text{salary} < 75000}(\text{instructor}))$$

But the number of rows involved in the projection operation may be (significantly) smaller in the second expression



Basic Steps in Query Processing

- Optimization
 - A relational algebra expression may have many equivalent expressions
 - ... and each relational algebra operation can be evaluated using one of several different algorithms
 - *Correspondingly, a relational-algebra expression can be evaluated in many ways*

Basic Steps in Query Processing

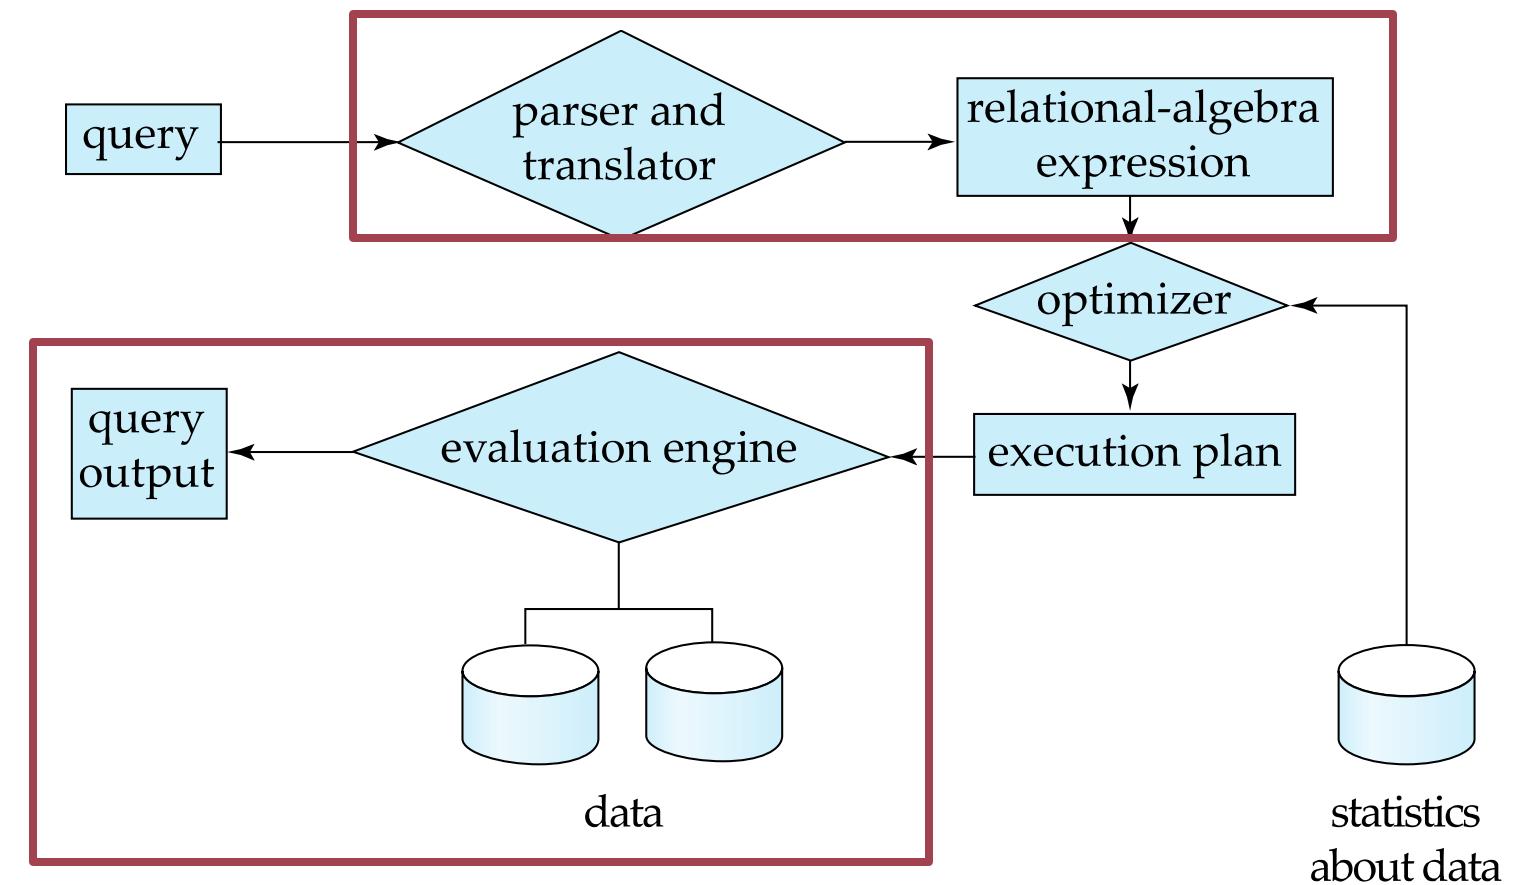
- Optimization
 - **Evaluation Plan:** Annotated expression specifying detailed evaluation strategy
 - E.g.,:
 - Use an index on salary to find instructors with $\text{salary} < 75000$
 - Or perform complete relation scan and discard instructors with $\text{salary} < 75000$

Query Optimization: Choose the one with the lowest cost among all equivalent evaluation plans

- Cost can be estimated using statistical information from the database catalog
 - E.g., Number of tuples in each relation, size of tuples, etc.

Basic Steps in Query Processing

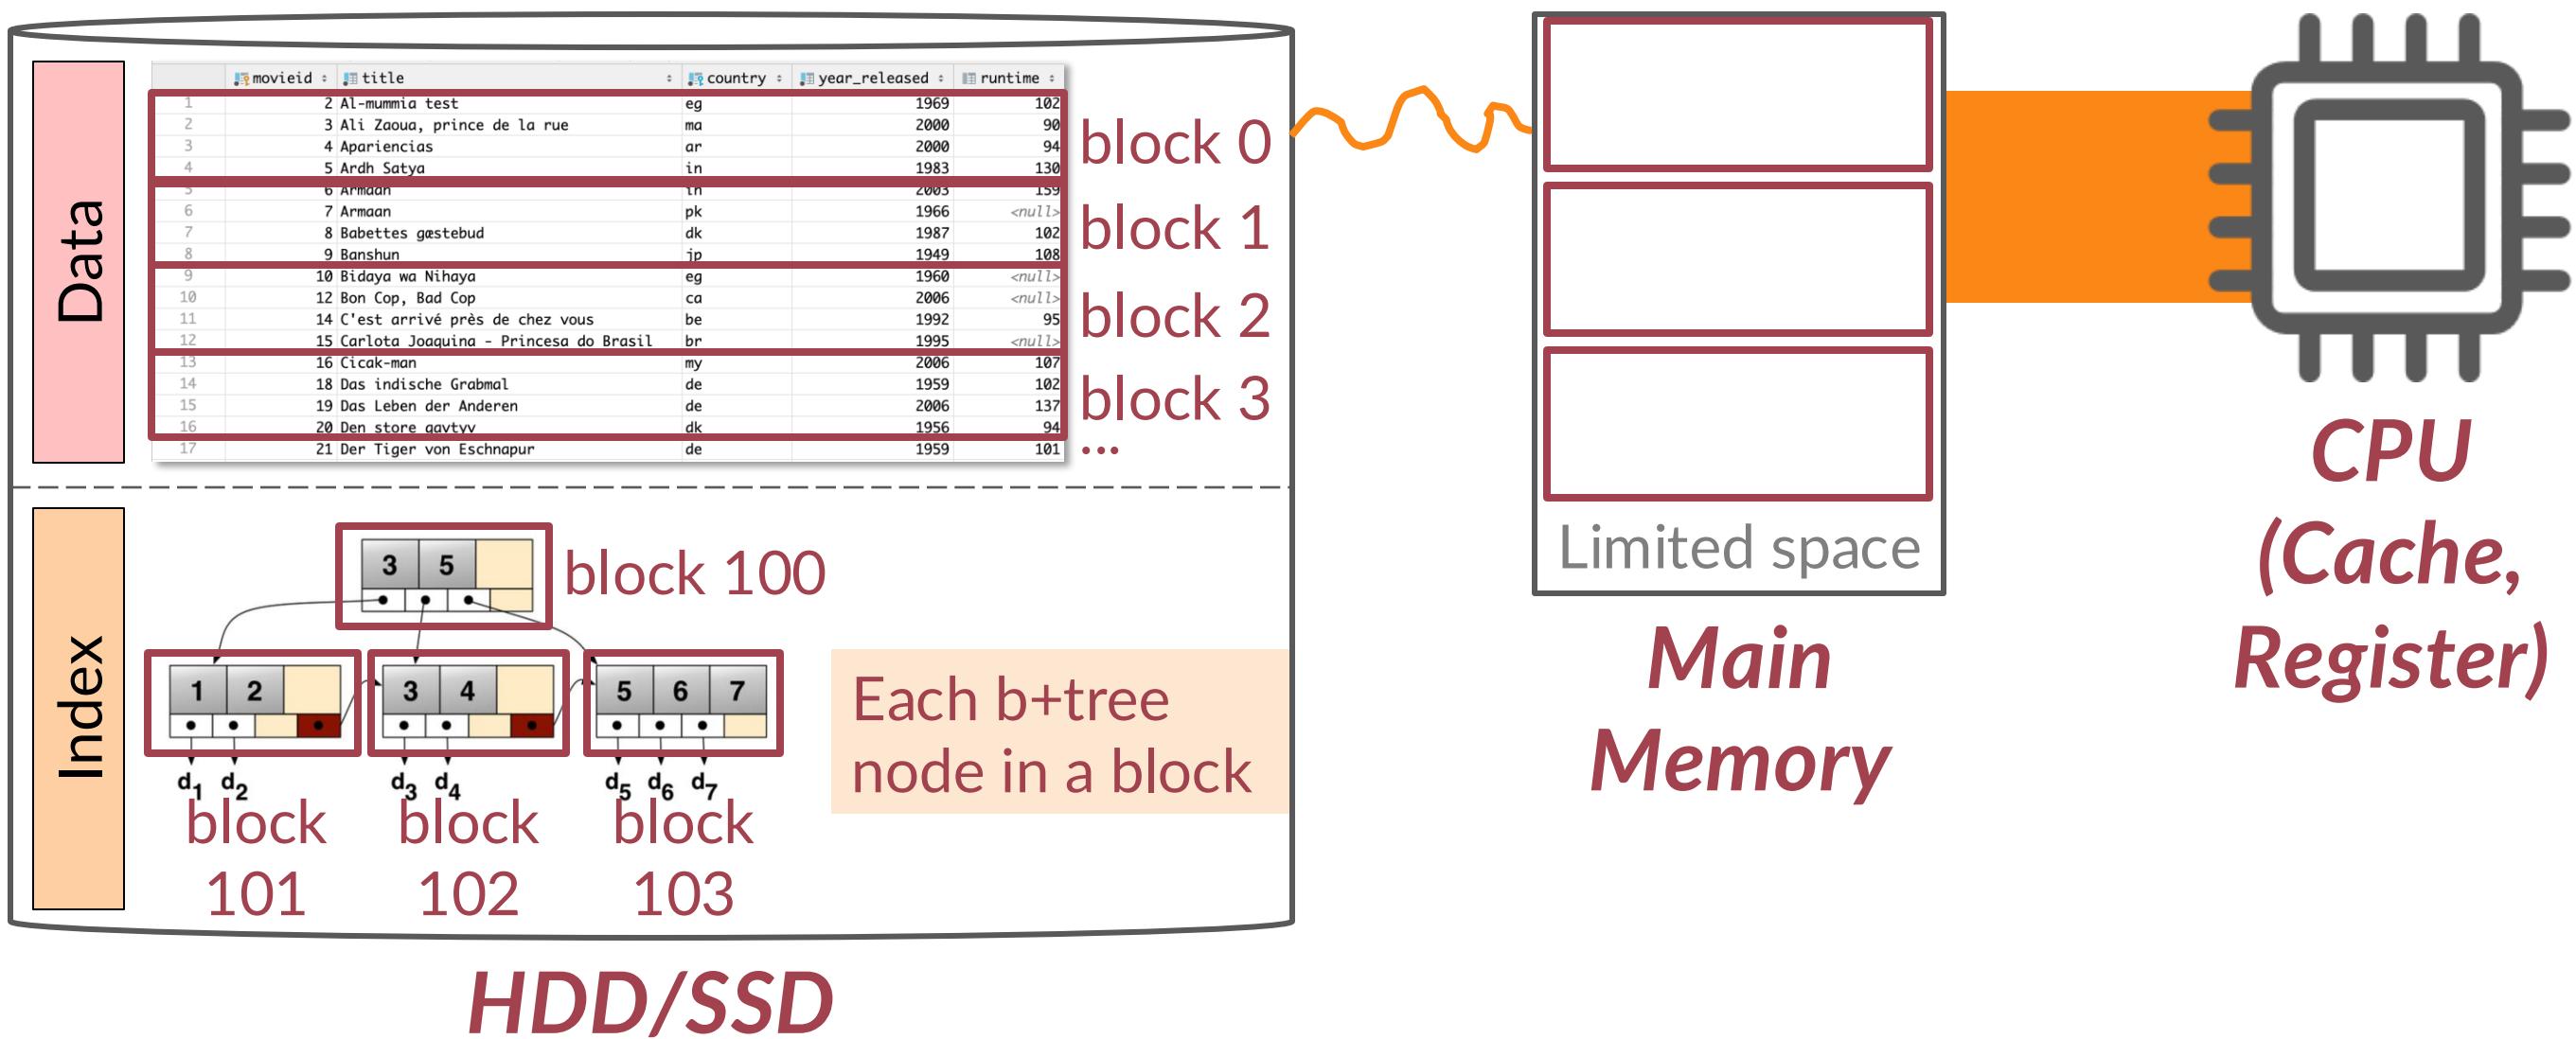
- Evaluation
 - The query-execution engine takes a query-evaluation plan, executes that plan, and returns the answers to the query



Cost of Query

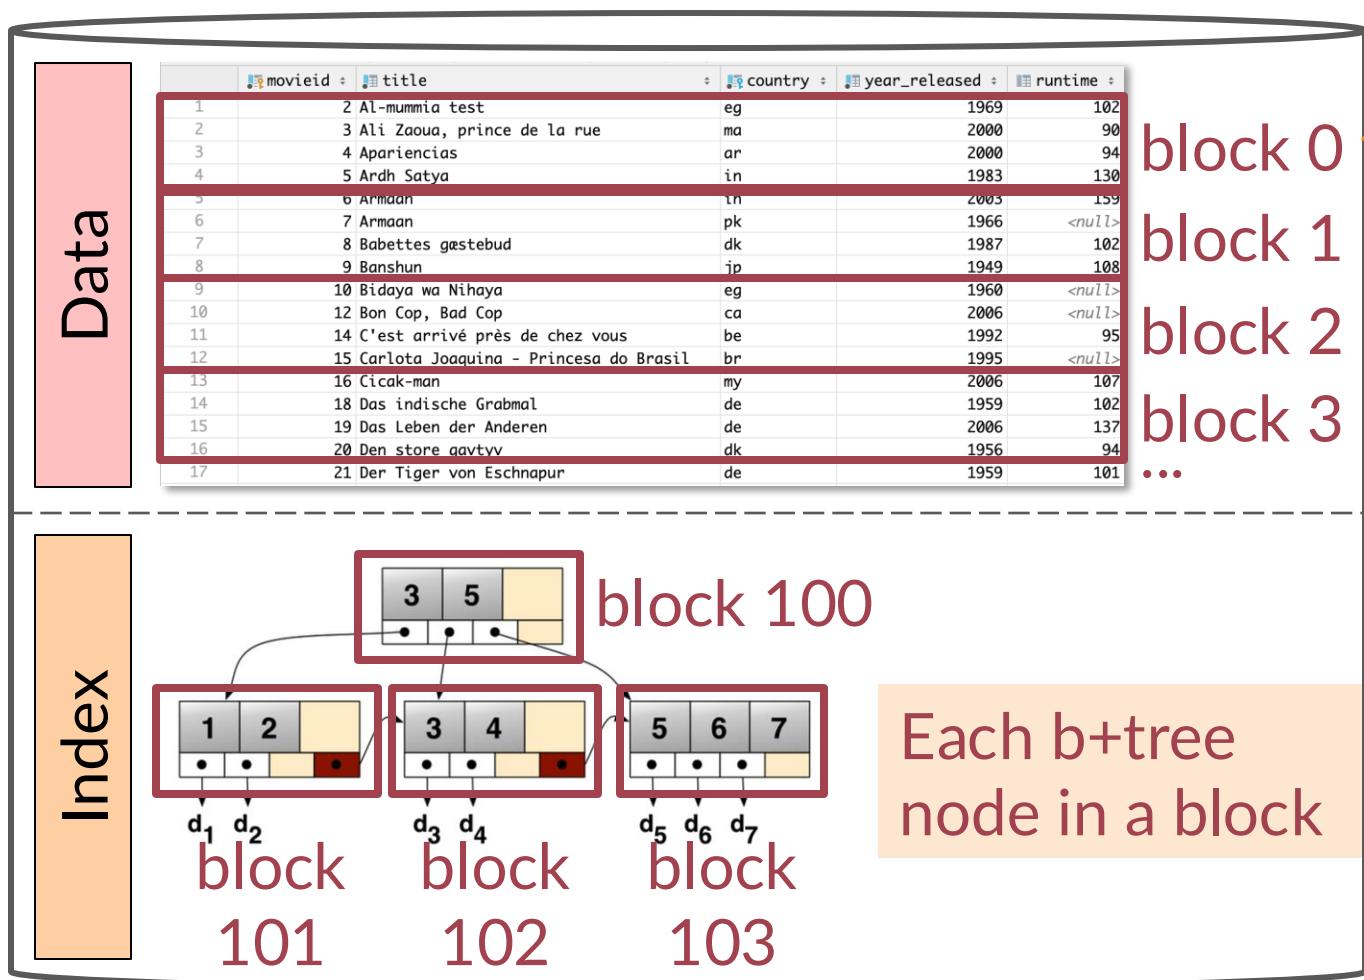
Prerequisite

- Storage model



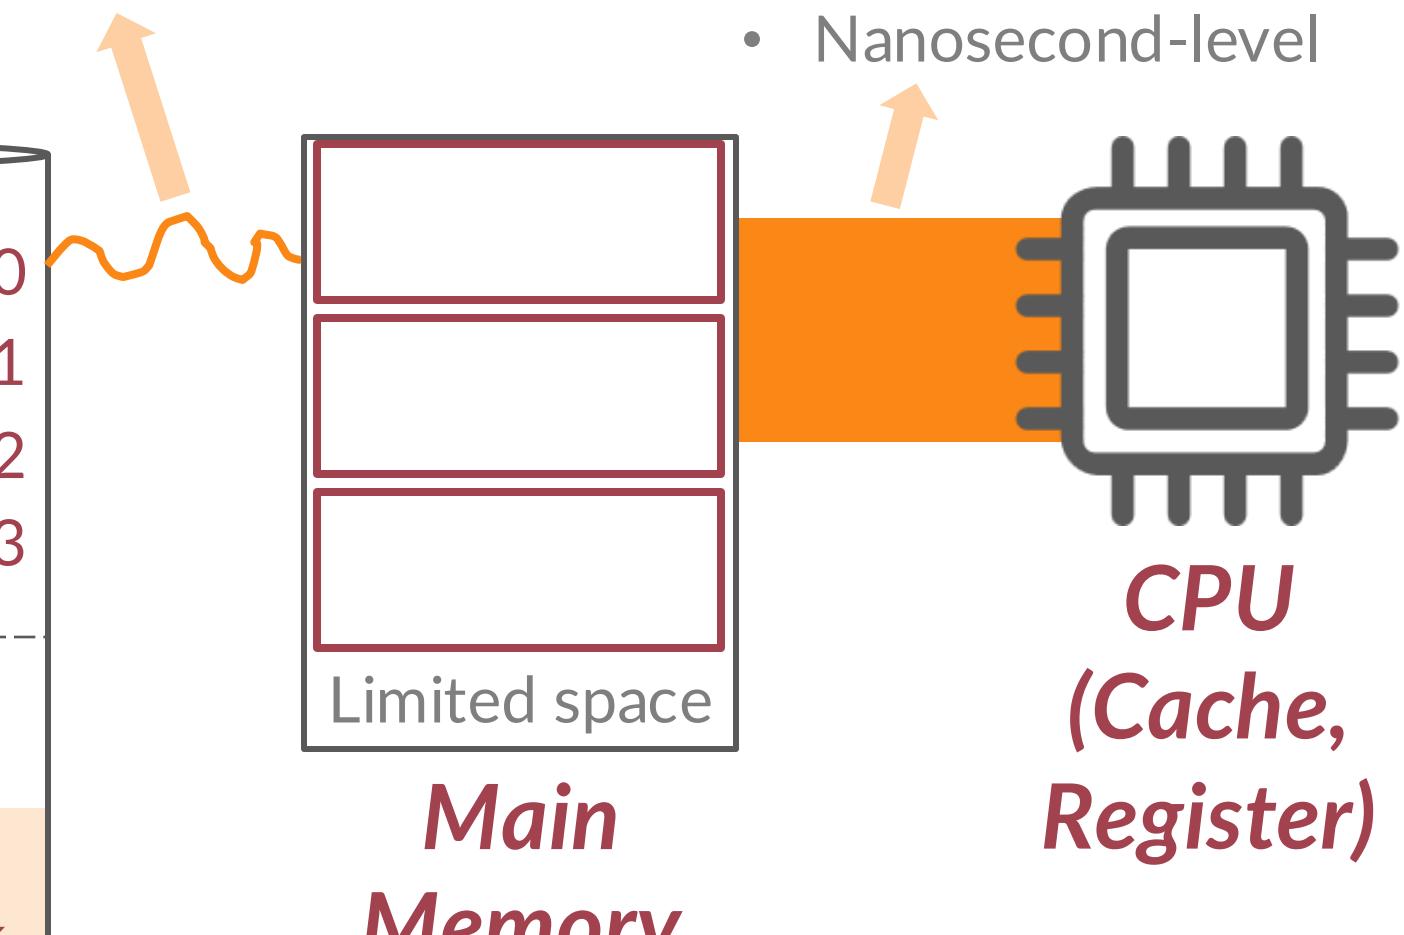
Prerequisite

- Storage model



HDD/SSD

- Relatively small bandwidth
 - 100MB/s ~ <10GB/s
- High latency
 - Millisecond-level
- Very large bandwidth
 - 94GB/s (for DDR4 2933*)
- Very low latency
 - Nanosecond-level



*

<https://www.intel.com/content/www/us/en/support/articles/000056722/processors/intel-core-processors.html>

Prerequisite

- Measuring query cost
 - Disk cost can be estimated as:
 - Number of seeks * average-seek-cost
 - Number of blocks read * average-block-read-cost
 - Number of blocks written * average-block-write-cost
 - For simplicity, we just use the **number of block transfers from disk and the number of seeks** as the cost measures
 - t_T – time to transfer one block
 - Assuming for simplicity that write cost is same as read cost
 - t_S – time for one seek
 - E.g., cost for b block transfers plus S seeks
$$b * t_T + S * t_S$$

Prerequisite

- Measuring query cost
 - t_S and t_T depend on where data is stored. With 4 KB blocks:
 - High end magnetic disk: $t_S = 4$ msec and $t_T = 0.1$ msec
 - SSD: $t_S = 20\text{-}90$ microsec and $t_T = 2\text{-}10$ microsec for 4KB
 - Required data may be buffer resident already, avoiding disk I/O
 - But hard to take into account for cost estimation
 - Worst case estimates assume that no data is initially in buffer and only the minimum amount of memory needed for the operation is available
 - But more optimistic estimates are used in practice
 - We ignore CPU costs for simplicity
 - Real systems do take CPU cost into account
 - Network costs must be considered for parallel systems

Overview

- Selection
- Joining

Selection Operation

- Let's start from this simple query:



```
select * from movies where [CONDITION];
```

- If you are the designer of the database engine, what do you think is the best way to fulfill this requirement?
- Two factors to consider:
 - What comparison is it in the CONDITION (equality / comparison)?
 - Does the column involved in the CONDITION have an index?

Basic Linear Scan

- Linear Search (displayed as Seq_Scan in PostgreSQL)
 - Scan each file block and test all records to see whether they satisfy the selection condition
 - Cost estimate = b_r block transfers + 1 seek
 - Assuming blocks of the file are stored contiguously
 - b_r denotes number of blocks containing records from relation r
- Although slower than other algorithms for implementing selection, linear search can be applied regardless of
 - Selection condition
 - Ordering of records in the file
 - Availability of indexes

Basic Linear Scan

- However, a full-table linear scan on extremely-large tables can be a disaster
 - E.g., billions of records in database
 - That's why we need other optimized ways

Index Scan

- **Index scan** – Search algorithms that use an index
 - Selection condition must be on search-key of index



```
select * from movies where movieid = 125;
```

We have a B+ tree index
on `movieid`

- Plan: Index Scan



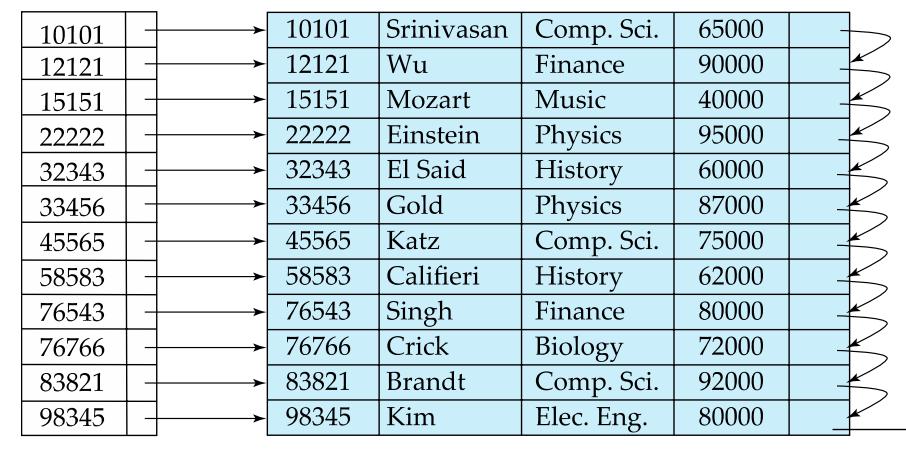
```
select * from movies where runtime = 100;
```

We don't have any index
on `runtime`

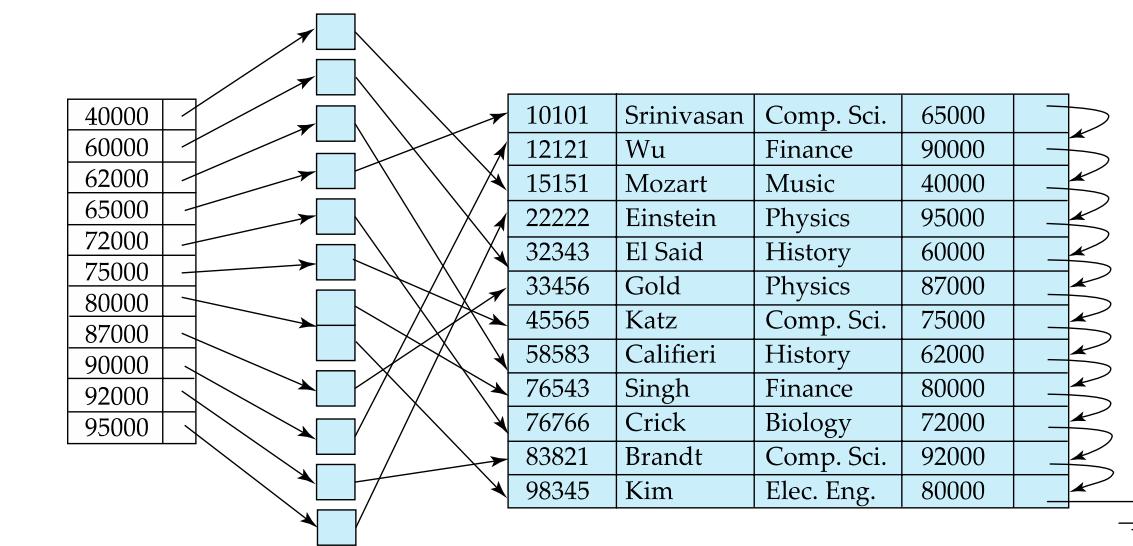
- Plan: Seq Scan

Index Scan

- **Index scan** – Search algorithms that use an index
 - Selection condition must be on search-key of index
- Unlike linear scan, we need to talk about different types of indexes and CONDITIONS
 - Clustered / Non-clustered index (Primary / Secondary index)
 - Equality / Comparison test



Clustered index



Non-clustered index

Index Scan

h_i : height of the B+-tree

Clustered index, equality on key

- Retrieve a single record that satisfies the corresponding equality condition
 - *key => no duplicated values*
 - $\text{Cost} = (h_i + 1) * (t_T + t_S)$
 - Index lookup traverses the height of the tree plus one I/O to fetch the record; each of these I/O operations requires a seek and a block transfer.

Clustered index, equality on non-key

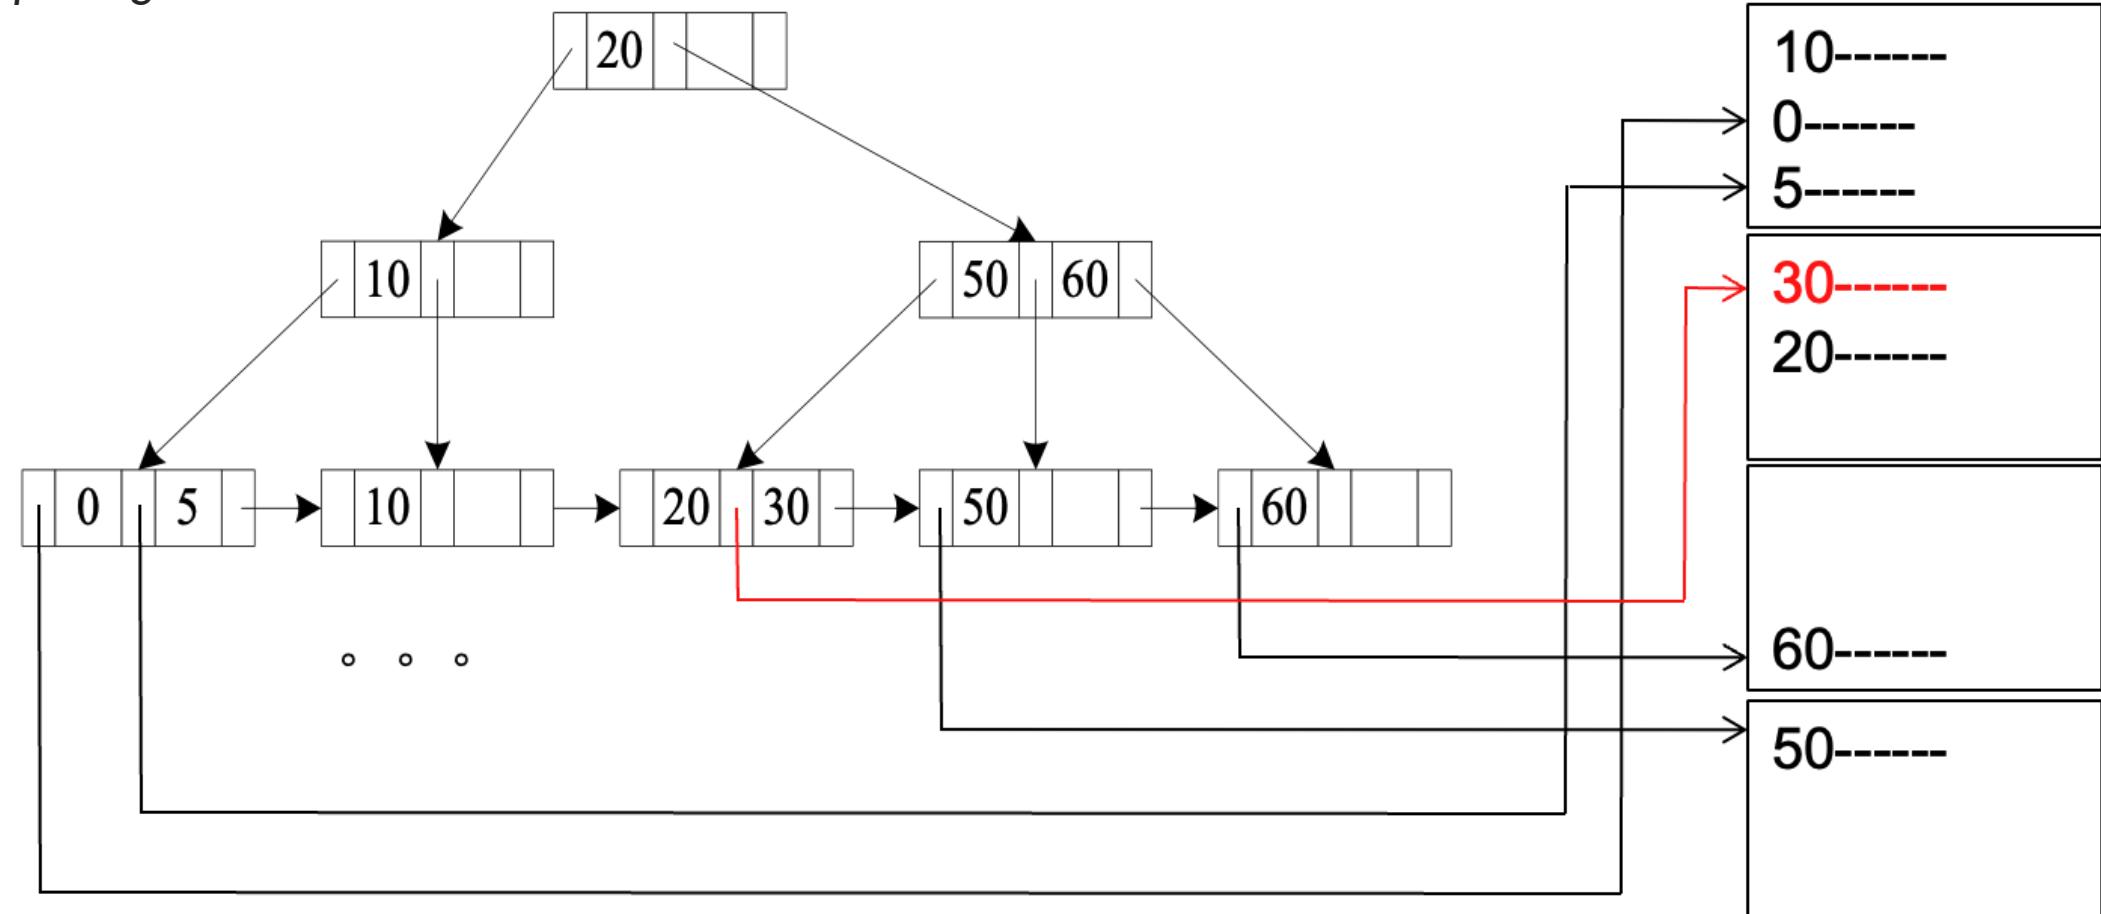
- Retrieve multiple records
 - *non key attributes => possible to have duplicated values*
 - $\text{Cost} = h_i * (t_T + t_S) + t_S + t_T * b$
 - One seek for each level of the tree, one seek for the first block
 - Let b = number of blocks containing matching records, which will be on consecutive blocks (since it is a clustering index) and don't require additional seeks

Index Scan

h_i : height of the B+-tree

Secondary index, equality on key/non-key

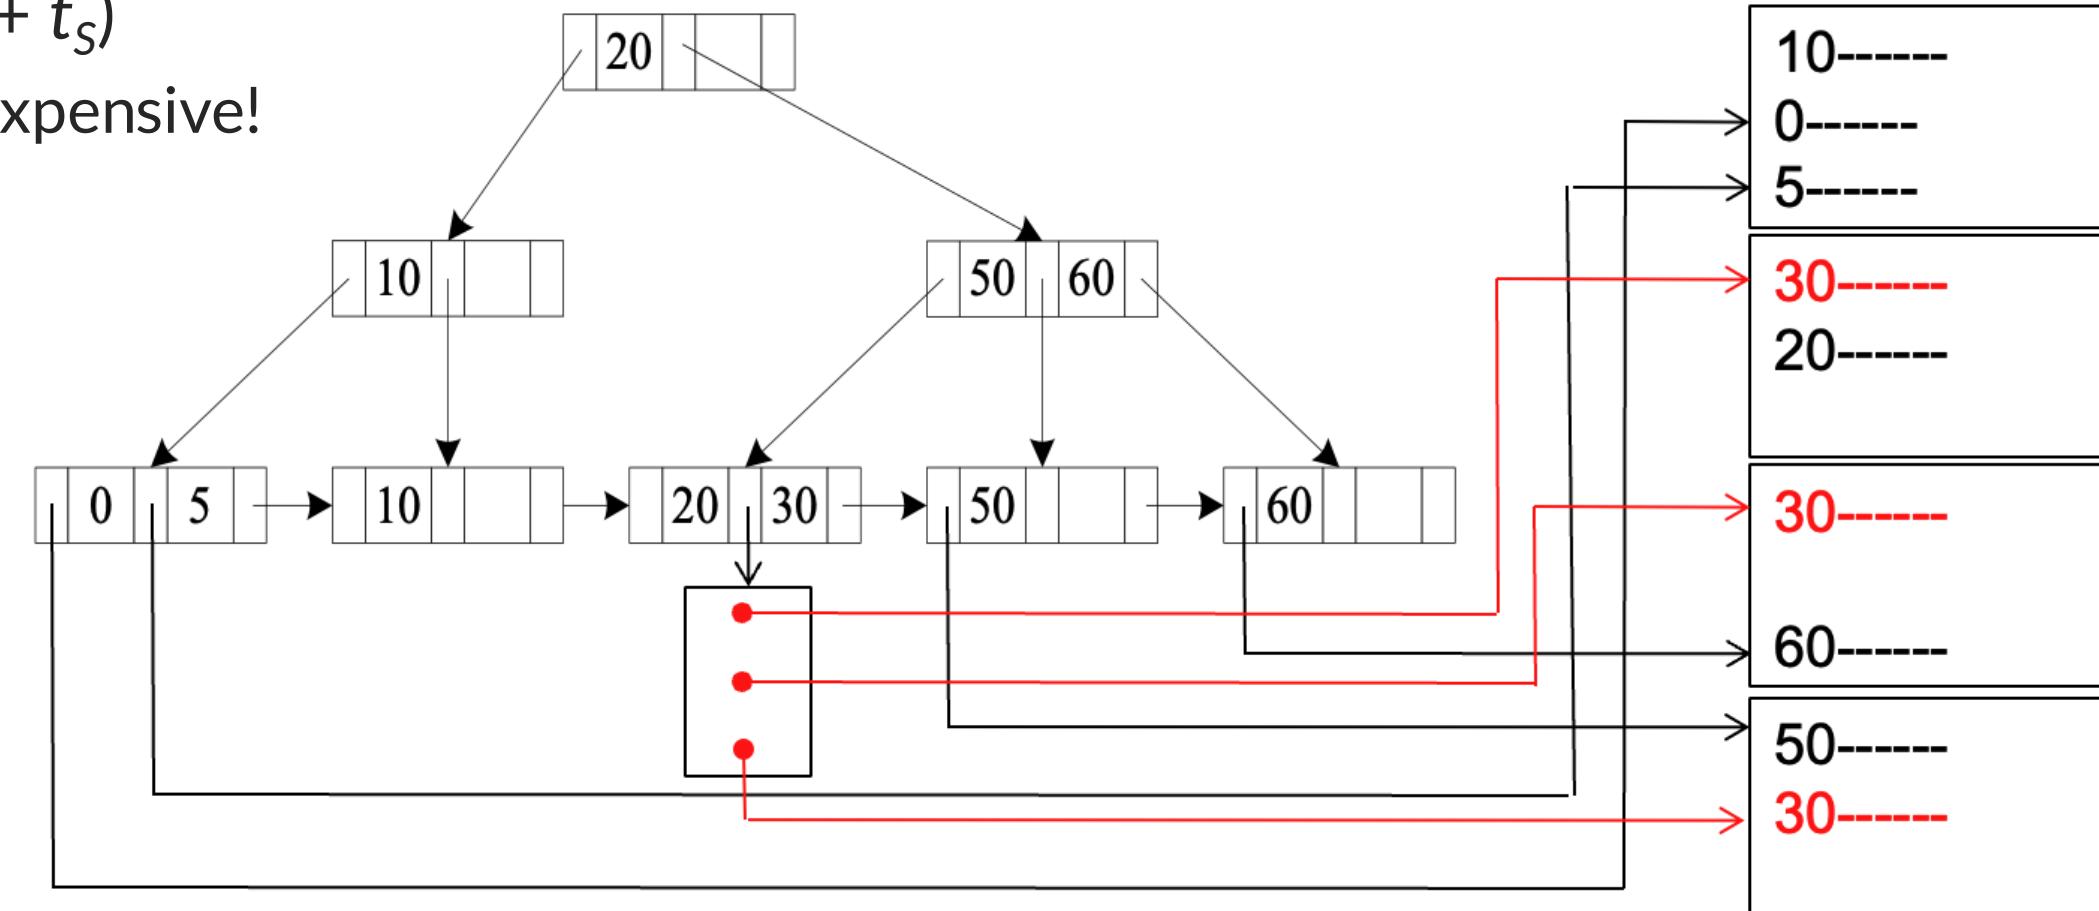
- Retrieve a single record if the search-key is a candidate key
 - $Cost = (h_i + 1) * (t_T + t_S)$



Index Scan

Secondary index, equality on key/non-key

- Retrieve multiple records if search-key is not a candidate key
 - Each of n matching records may be on a different block
 - Cost = $(h_i + n) * (t_T + t_S)$
 - Can be very expensive!

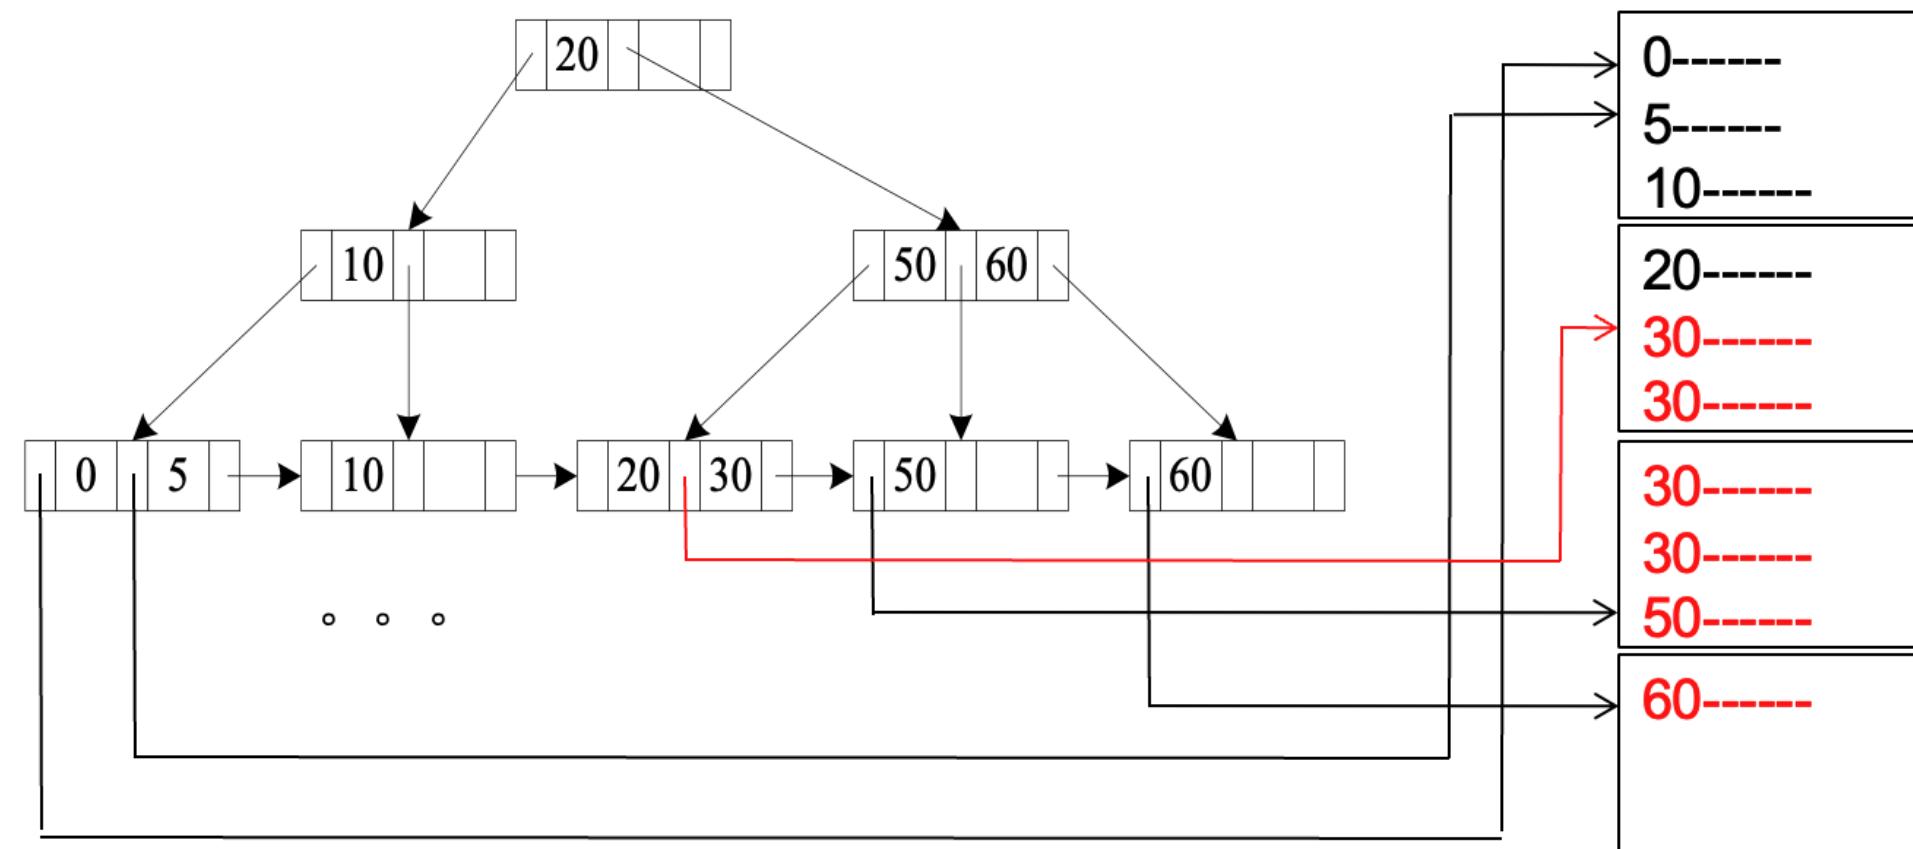


Index Scan

Tip: Comparison tests can always be fulfilled with linear scans, which is the fallback solution

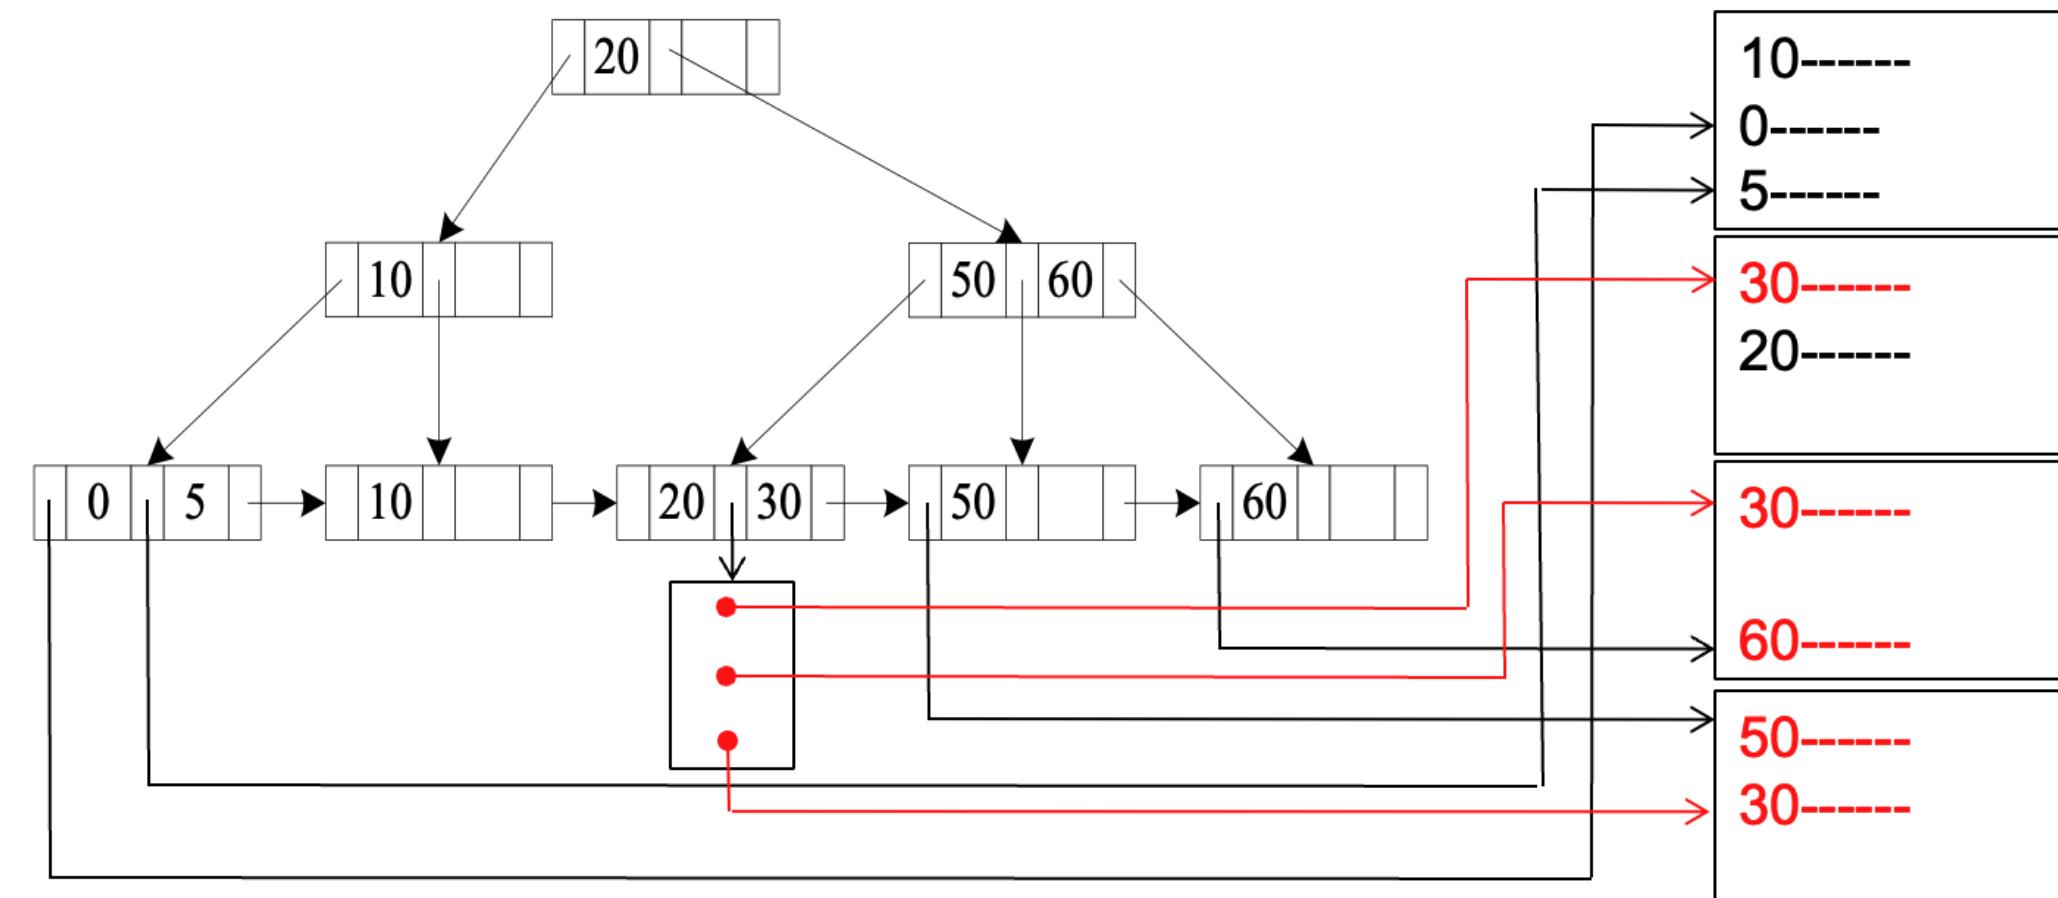
- Clustered index, comparison (i.e., Relation is sorted on A)

- For $\sigma_{A \geq v}(r)$, use index to find first tuple $\geq v$ and scan relation sequentially from there
- For $\sigma_{A \leq v}(r)$, just scan relation sequentially till first tuple $> v$; do not use index



Index Scan

- Non-clustered index, comparison
 - For $\sigma_{A \geq v}(r)$, use index to find first index entry $\geq v$ and scan index sequentially from there, to find pointers to records.
 - For $\sigma_{A \leq v}(r)$, just scan leaf pages of index finding pointers to records, till first entry $> v$
- In either case, retrieving records that are pointed to requires an I/O per record
- Linear scan may be cheaper!



Complex Selections

Conjunction: $\sigma_{\theta_1} \wedge \theta_2 \wedge \dots \wedge \theta_n(r)$

- **Conjunctive selection using single-key index(es)**

- Select a θ_i and algorithms mentioned above that results in the least cost for $\sigma_{\theta_i}(r)$
- Test other conditions on tuple after fetching it into memory buffer

- **Conjunctive selection using multi-key index**

- Use appropriate composite (multiple-key) index if available

- **Conjunctive selection by intersection of identifiers**

- Requires indices with record pointers/identifiers
 - Here, “indices” means the order number of records in the files, like array indices. They are not indexes
- Use corresponding index for each condition, and take intersection of all the obtained sets of record pointers
- Then fetch records from file (and test remaining conditions)

Complex Selections

Disjunction: $\sigma_{\theta_1} \vee \theta_2 \vee \dots \vee \theta_n(r)$

Disjunctive selection by union of identifiers

(Similar to the third way on the previous page)

- Applicable if *all* conditions have available indexes
 - Use corresponding index for each condition, and take union of all the obtained sets of record pointers
 - Then fetch records from file
- Otherwise, just use linear scan
 - The disjunctive condition tested on each tuple during the scan

How join Works (with the help of indexes)

- Some widely-used join algorithms
 - Nested-loop join
 - Indexed nested-loop join
 - Merge join

Nested-loop Join

- To compute the *theta* join $r \bowtie_{\theta} s$

```
for each tuple  $t_r$  in  $r$  do begin  
    for each tuple  $t_s$  in  $s$  do begin  
        test pair  $(t_r, t_s)$  to see if they satisfy the join condition  $\theta$   
        if they do, add  $t_r \bullet t_s$  to the result.  
    end  
end
```
- r is called the **outer relation** and s the **inner relation** of the join
 - Think about the “outer loop” and the “inner loop” in programming
- Requires no indices and can be used with any kind of join condition
 - Expensive since it examines every pair of tuples in the two relations

Nested-loop Join via File Scan

- In the worst case, if the memory can only hold one block of each relation, the estimated cost is:
 - $n_r * b_s + b_r$ block transfers, plus $n_r + b_r$ seeks
 - n_r - number of records in relation r
 - b_s, b_r - number of blocks in relation s and r
 - $n_r * b_s + b_r$: for each record in r , need to read all blocks in s ; and need to read all b_r blocks in r
 - $n_r + b_r$: for each record in r , need to do seek once for s ; and need to seek r for b_r times
 - Assuming r and s are stored contiguously on disks
- If the smaller relation fits entirely in memory, use that as the inner relation
 - Reduces cost to $b_r + b_s$ block transfers and 1 + 1 seeks

Indexed Nested-Loop Join

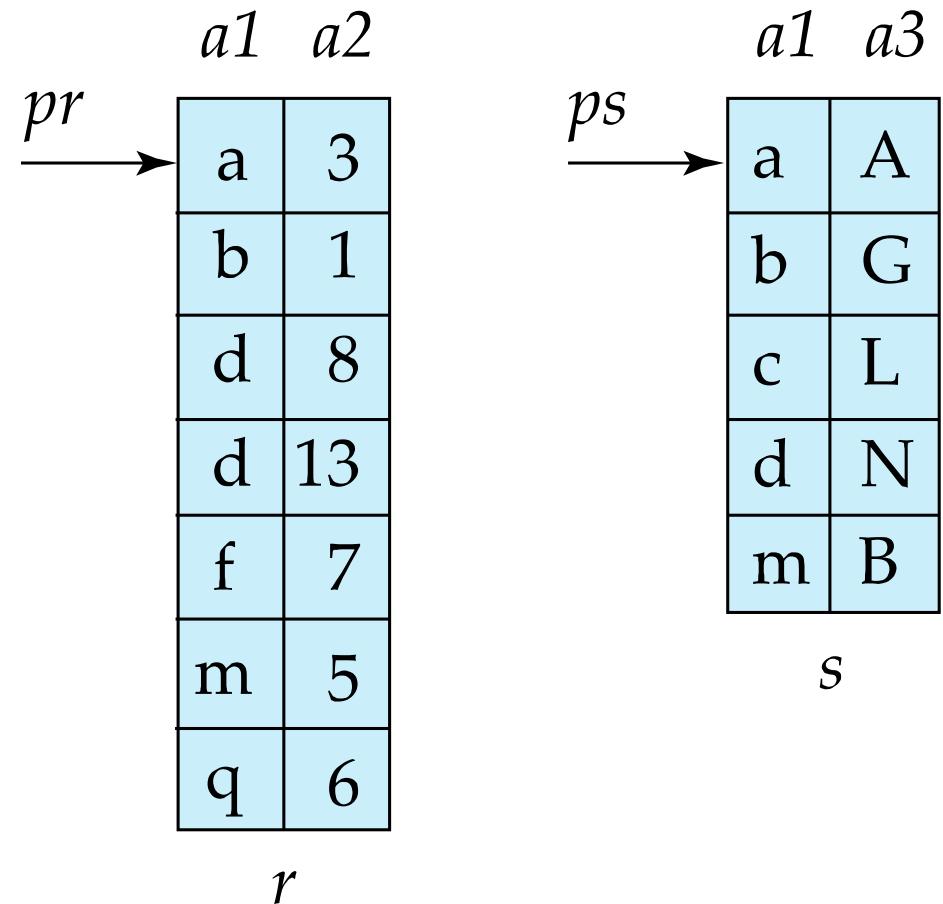
- Index lookups can replace file scans if
 - join is an equi-join ($r \bowtie_{r.A=s.B} s$) or natural join and
 - an index is available on the inner relation's join attribute
 - Can construct an index just to compute a join
- For each tuple t_r in the outer relation r , use the index to look up tuples in s that satisfy the join condition with tuple t_r
 - Essentially a selection operation given the values of the joining attribute in t_r

Indexed Nested-Loop Join

- Worst case: Buffer in memory has space for only one page of r , and, for each tuple in r , we perform an index lookup on s
- Cost of the join: $b_r * (t_T + t_S) + n_r * c$
 - Where c is the cost of traversing index and fetching all matching s tuples for one tuple of r
 - c can be estimated as cost of a single selection on s using the join condition
 - We need b_r seeks as the disk head may have moved between each I/O
- If indices are available on join attributes of both r and s , use the relation with fewer tuples as the outer relation

Merge Join

- a.k.a., sort-merge join
 - Zipper-like joining
- Steps
 - Sort both relations on their join attribute (if not already sorted on the join attributes)
 - Merge the sorted relations to join them
 - if $r.a1[pr] < s.a1[ps]$, $pr++$
 - elif $r.a1[pr] > s.a1[ps]$, $ps++$
 - else, join and move pr and ps
- Join step is similar to the merge stage of the sort-merge algorithm
 - Main difference is handling of duplicate values in join attribute – every pair with same value on join attribute must be matched



Merge Join

- Can be used only for equi-joins ($r \bowtie_{r.A=s.B} s$) and natural joins
- Once the relations are in sorted order, each block needs to be read only once
 - Assuming all tuples for any given value of the join attributes fit in memory
- The cost of merge join is:
 $b_r + b_s$ block transfers + $\lceil b_r / b_b \rceil + \lceil b_s / b_b \rceil$ seeks
+ the cost of sorting if relations are unsorted

b_b : memory buffer size, counted in number of blocks, for each relation

Hash Join

- Hash join
 - Build a set of buckets for a smaller table to speed up the data lookup
- Procedure:
 - 1. Create a hash table for the smaller table t_1 in the memory
 - 2. Scan the larger table t_2 . For each record r ,
 - 2.1 Compute the hash value of $r.\text{join_attribute}$
 - 2.2 Map to corresponding rows in t_1 using the hash table

Principles of Database Systems (CS307)

Lecture 14 part 2: Query Optimization

Zhong-Qiu Wang

Department of Computer Science and Engineering
Southern University of Science and Technology

- Most contents are from slides made by Stéphane Faroult and the authors of Database System Concepts (7th Edition).
- Their original slides have been modified to adapt to the schedule of CS307 at SUSTech.
- The slides are largely based on the slides provided by Dr. Yuxin Ma

Query Optimization

- Purpose of query optimization
 - Select an effective way to retrieve the data based on queries while spending the least computational effort
 - However, it is only “spending less computational effort” in most scenarios, not least

Query Optimization

- Users don't need to consider the best way of writing queries
 - We want DBMS to construct a query-evaluation plan that minimizes the cost of query evaluation
- Automated optimization can perform better (for most of the time)
 - Utilize the data dictionary
 - Real-time utilization based on physical storage changes
 - Optimizer can evaluate hundreds of execution plans in a very short time compared with human programmers
 - Human users do not need to learn advanced optimization techniques any more, which is conducted by optimizers instead

An Example in the Movie Dataset

- The same query can be represented in different plans
 - E.g., retrieve the titles of those movies from China



```
select m.title  
from movies m, countries c  
where m.country = c.country_code and c.country_name = 'China';
```

An Example in the Movie Dataset

- The corresponding relational algebra expressions:

(1) $\Pi_{title} (\sigma_{movies.country = countries.country_code \wedge countries.country = "China"} (movies \times countries))$

(2) $\Pi_{title} (\sigma_{countries.country = "China"} (movies \bowtie_{movies.country = countries.country_code} countries))$

(3) $\Pi_{title} (movies \bowtie_{movies.country = countries.country_code} (\sigma_{countries.country = "China"} (countries)))$

An Example in the Movie Dataset

- The corresponding relational algebra expressions:

(1) $\prod_{title} (\sigma_{movies.country = countries.country_code \wedge countries.country = "China"} (movies \times countries))$

(2) $\prod_{title} (\sigma_{countries.country = "China"} (movies \bowtie_{movies.country = countries.country_code} countries))$

(3) $\prod_{title} (movies \bowtie_{movies.country = countries.country_code} (\sigma_{countries.country = "China"} (countries)))$

- In (1), **a full Cartesian product** will be computed, which costs huge time for matching all pairs and massive temporary storage space for the intermediate product table
- In (2), **a smaller intermediate join table** is to be cached, which saves some space
- In (3), **the filter** ($\sigma_{c.country = "China"}$) **reduces the size of the right table** in the join operation, which saves a lot of time for pair matching and caching intermediate join table

An Example in the Movie Dataset

- In addition, the filter operation can be further accelerated once an index is built upon the *country* column

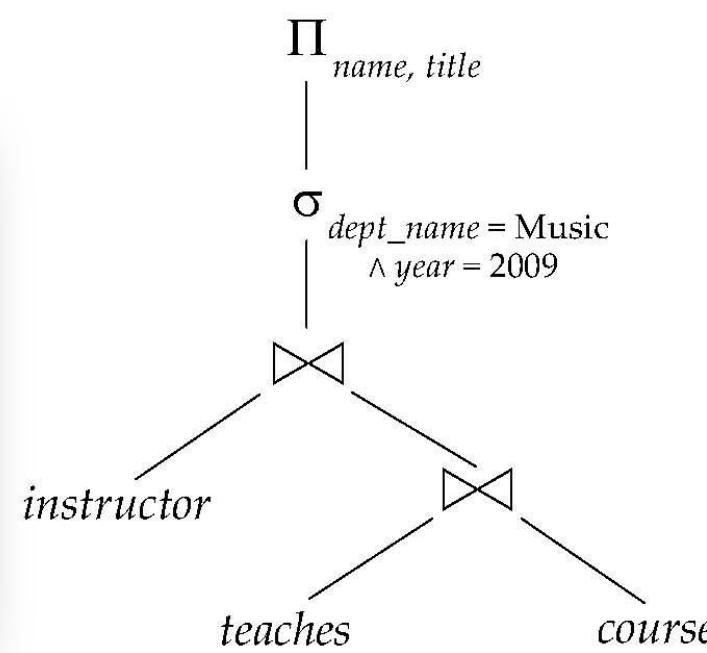
$$(3) \Pi_{title} (\text{movies} \bowtie_{\text{movies.country} = \text{countries.country_code}} (\sigma_{\text{countries.country} = "China"}(\text{countries}))$$

Generating Equivalent Expressions

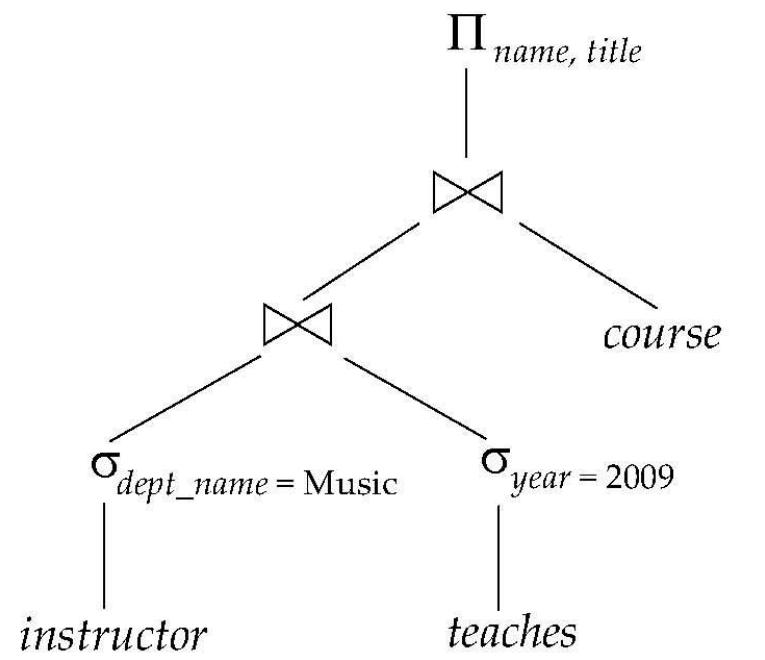
- Alternative ways of evaluating a given query
 - Equivalent expressions
 - Different algorithms for each operation



```
select name, title  
from instructor  
    natural join (teaches natural join course)  
where dept_name = 'Music' and year = 2009;
```



(a) Initial expression tree



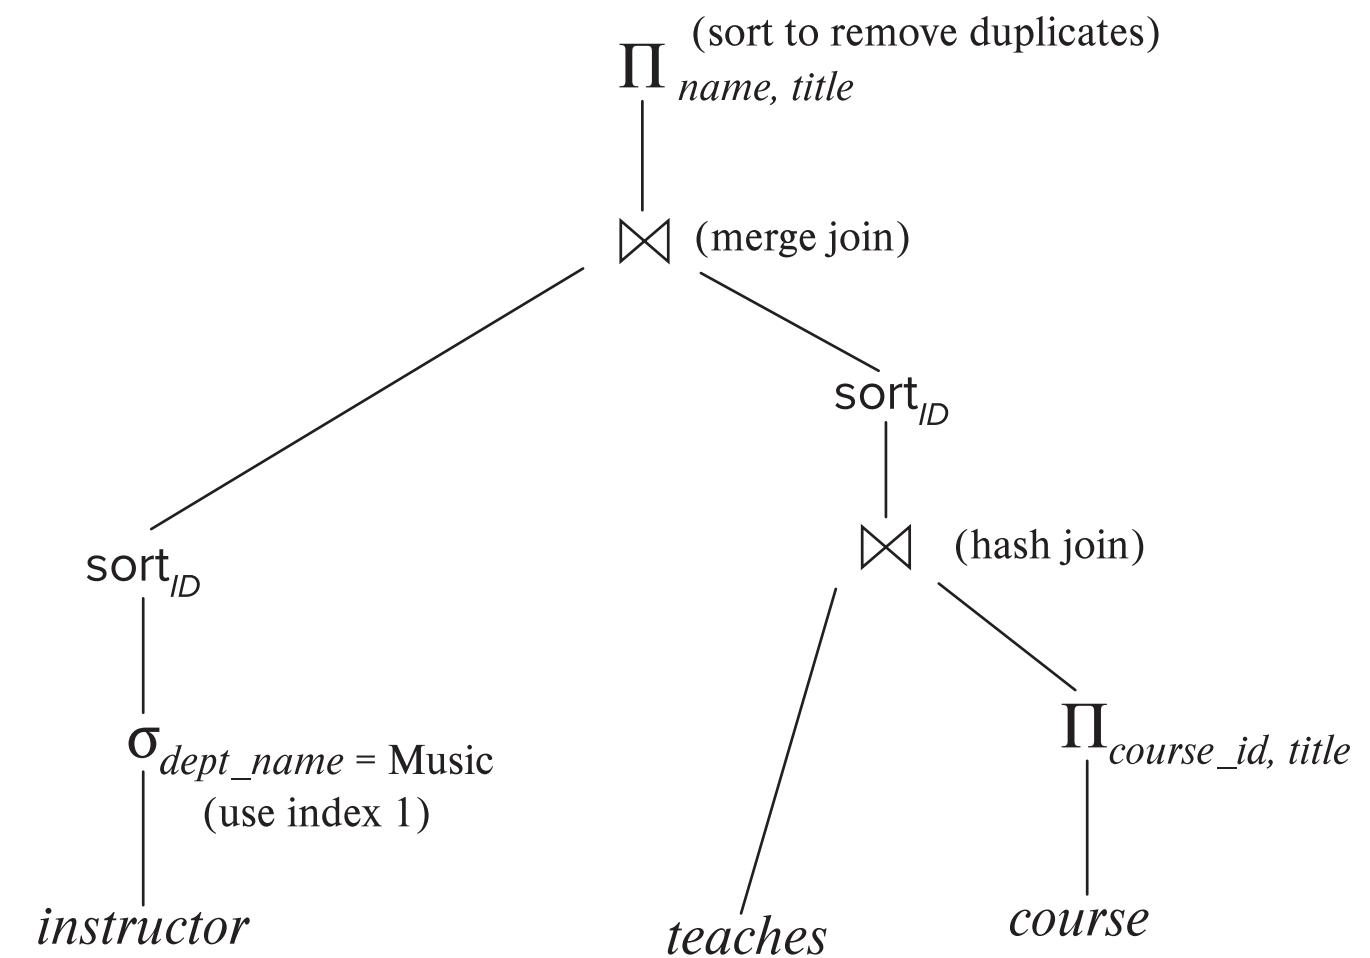
(b) Tree after multiple transformations

Generating Equivalent Expressions

- An **evaluation plan** defines exactly what algorithm is used for each operation, and how the execution of the operations is coordinated



```
select name, title  
from instructor  
    natural join (teaches natural join course)  
where dept_name = 'Music' and year = 2009;
```



Transformation of Relational Expressions

- Two relational algebra expressions are said to be **equivalent** if the two expressions generate the same set of tuples on every legal database instance
 - Note: order of tuples is irrelevant
 - We don't care if they generate different results on databases that violate integrity constraints
- An **equivalence rule** says that expressions of two forms are equivalent
 - ... i.e., we can replace expression of the first form by second, or vice versa
 - The optimizer uses equivalence rules to transform expressions into other logically equivalent expressions

Equivalence Rules

1. Conjunctive selection operations can be deconstructed into a sequence of individual selections.

$$\sigma_{\theta_1 \wedge \theta_2}(E) \equiv \sigma_{\theta_1}(\sigma_{\theta_2}(E))$$

2. Selection operations are **commutative**

$$\sigma_{\theta_1}(\sigma_{\theta_2}(E)) \equiv \sigma_{\theta_2}(\sigma_{\theta_1}(E))$$

3. Only the last in a sequence of projection operations is needed, the others can be omitted

$$\Pi_{L_1}(\Pi_{L_2}(\dots(\Pi_{L_n}(E))\dots)) \equiv \Pi_{L_1}(E)$$

where $L_1 \subseteq L_2 \dots \subseteq L_n$

4. Selections can be combined with Cartesian products and theta joins

- $\sigma_\theta(E_1 \times E_2) \equiv E_1 \bowtie_\theta E_2$ (same as the definition of theta-join)
- $\sigma_{\theta_1}(E_1 \bowtie_{\theta_2} E_2) \equiv E_1 \bowtie_{\theta_1 \wedge \theta_2} E_2$

Equivalence Rules

5. Theta-join operations (and natural joins) are commutative.

$$E_1 \bowtie E_2 \equiv E_2 \bowtie E_1$$

6. (a) Natural join operations are associative:

$$(E_1 \bowtie E_2) \bowtie E_3 \equiv E_1 \bowtie (E_2 \bowtie E_3)$$

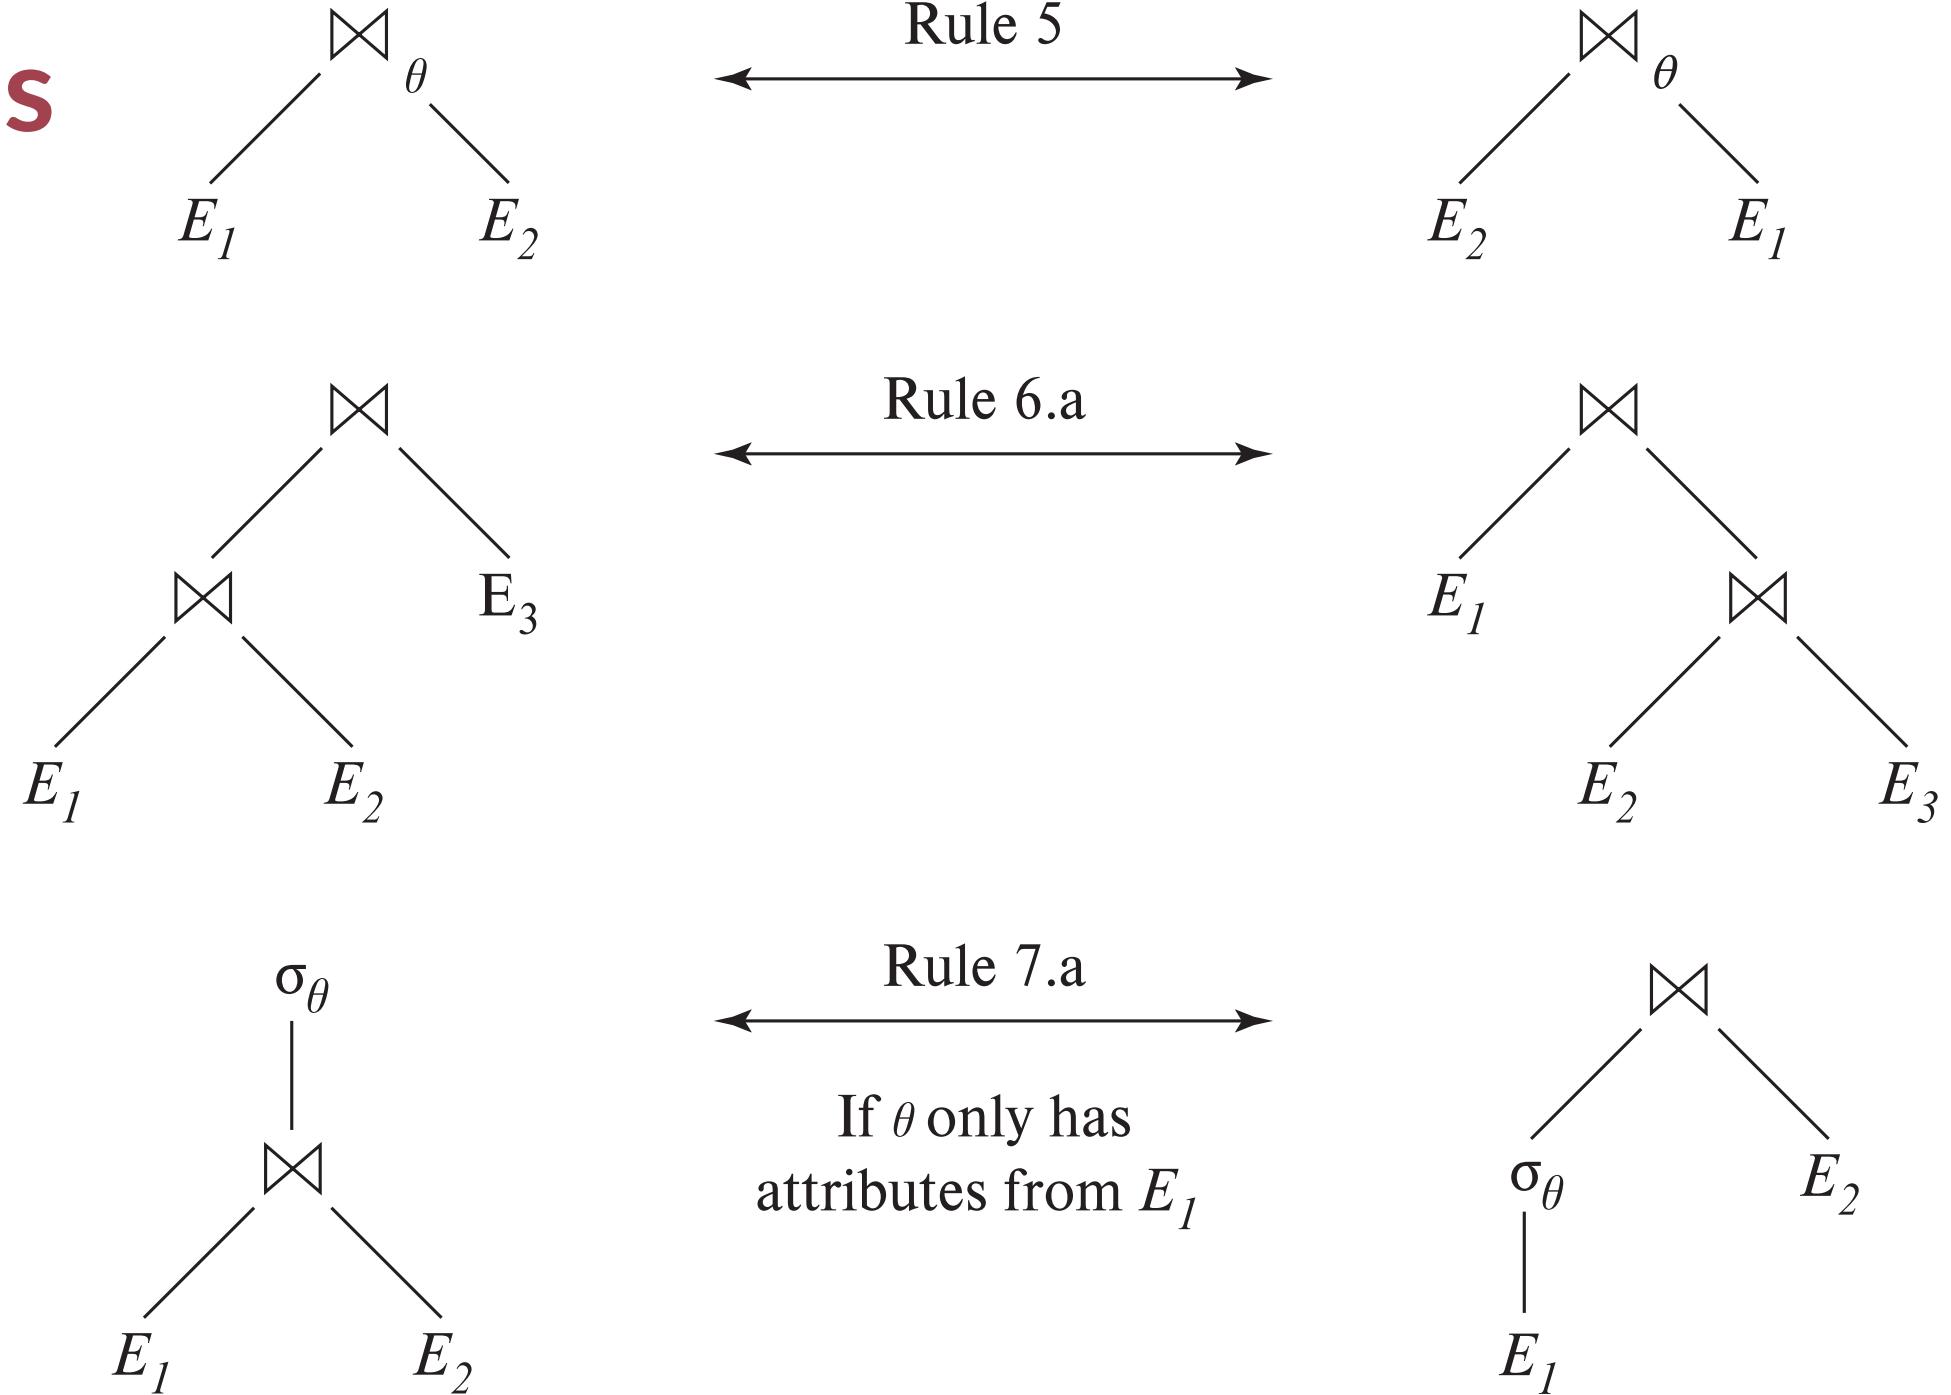
(b) Theta joins are associative in the following manner:

$$(E_1 \bowtie_{\theta_1} E_2) \bowtie_{\theta_2 \wedge \theta_3} E_3 \equiv E_1 \bowtie_{\theta_1 \wedge \theta_3} (E_2 \bowtie_{\theta_2} E_3)$$

where θ_2 involves attributes from only E_2 and E_3

Equivalence Rules

- Representation of Rule 5, 6(a) and 6(b) with diagrams



Equivalence Rules

7. The selection operation distributes over the theta join operation under the following two conditions:

- (a) When all the attributes in θ_0 involve only the attributes of one of the expressions (E_1) being joined:

$$\sigma_{\theta_0}(E_1 \bowtie_{\theta} E_2) \equiv (\sigma_{\theta_0}(E_1)) \bowtie_{\theta} E_2$$

- (b) When θ_1 involves only the attributes of E_1 and θ_2 involves only the attributes of E_2 :

$$\sigma_{\theta_1 \wedge \theta_2}(E_1 \bowtie_{\theta} E_2) \equiv (\sigma_{\theta_1}(E_1)) \bowtie_{\theta} (\sigma_{\theta_2}(E_2))$$

Perform selection early

Equivalence Rules

8. The projection operation distributes over the theta join operation as follows:

(a) If θ involves only attributes from $L_1 \cup L_2$:

- Let L_1 and L_2 be sets of attributes from E_1 and E_2 , respectively,

$$\Pi_{L_1 \cup L_2}(E_1 \bowtie_{\theta} E_2) \equiv \Pi_{L_1}(E_1) \bowtie_{\theta} \Pi_{L_2}(E_2)$$

(b) In general, consider a join $E_1 \bowtie_{\theta} E_2$:

- Let L_1 and L_2 be sets of attributes from E_1 and E_2 , respectively,
- Let L_3 be attributes of E_1 that are involved in join condition θ , but are not in L_1 , and,
- Let L_4 be attributes of E_2 that are involved in join condition θ , but are not in L_2 :

$$\Pi_{L_1 \cup L_2}(E_1 \bowtie_{\theta} E_2) \equiv \Pi_{L_1 \cup L_2}(\Pi_{L_1 \cup L_3}(E_1) \bowtie_{\theta} \Pi_{L_2 \cup L_4}(E_2))$$

* Similar equivalences hold for left, right, and full outer join operations: \bowtie , \bowtie_l , and \bowtie_r

Equivalence Rules

9. The set operations union and intersection are commutative

$$E_1 \cup E_2 \equiv E_2 \cup E_1$$

$$E_1 \cap E_2 \equiv E_2 \cap E_1$$

- However, set difference is not commutative

10. Set union and intersection are associative

$$(E_1 \cup E_2) \cup E_3 \equiv E_1 \cup (E_2 \cup E_3)$$

$$(E_1 \cap E_2) \cap E_3 \equiv E_1 \cap (E_2 \cap E_3)$$

Equivalence Rules

11. The selection operation distributes over \cup , \cap and $-$

- (a) $\sigma_{\theta}(E_1 \cup E_2) \equiv \sigma_{\theta}(E_1) \cup \sigma_{\theta}(E_2)$
- (b) $\sigma_{\theta}(E_1 \cap E_2) \equiv \sigma_{\theta}(E_1) \cap \sigma_{\theta}(E_2)$
- (c) $\sigma_{\theta}(E_1 - E_2) \equiv \sigma_{\theta}(E_1) - \sigma_{\theta}(E_2)$
- (d) $\sigma_{\theta}(E_1 \cap E_2) \equiv \sigma_{\theta}(E_1) \cap E_2$
- (e) $\sigma_{\theta}(E_1 - E_2) \equiv \sigma_{\theta}(E_1) - E_2$

12. The projection operation distributes over union

$$\Pi_L(E_1 \cup E_2) \equiv (\Pi_L(E_1)) \cup (\Pi_L(E_2))$$

Transformation Example: Pushing Selections

- Query: Find the names of all instructors in the Music department, along with the titles of the courses (in the Music department) that they teach

```
select name, title  
from instructor natural join (teaches natural join course  
where dept_name = 'Music';
```

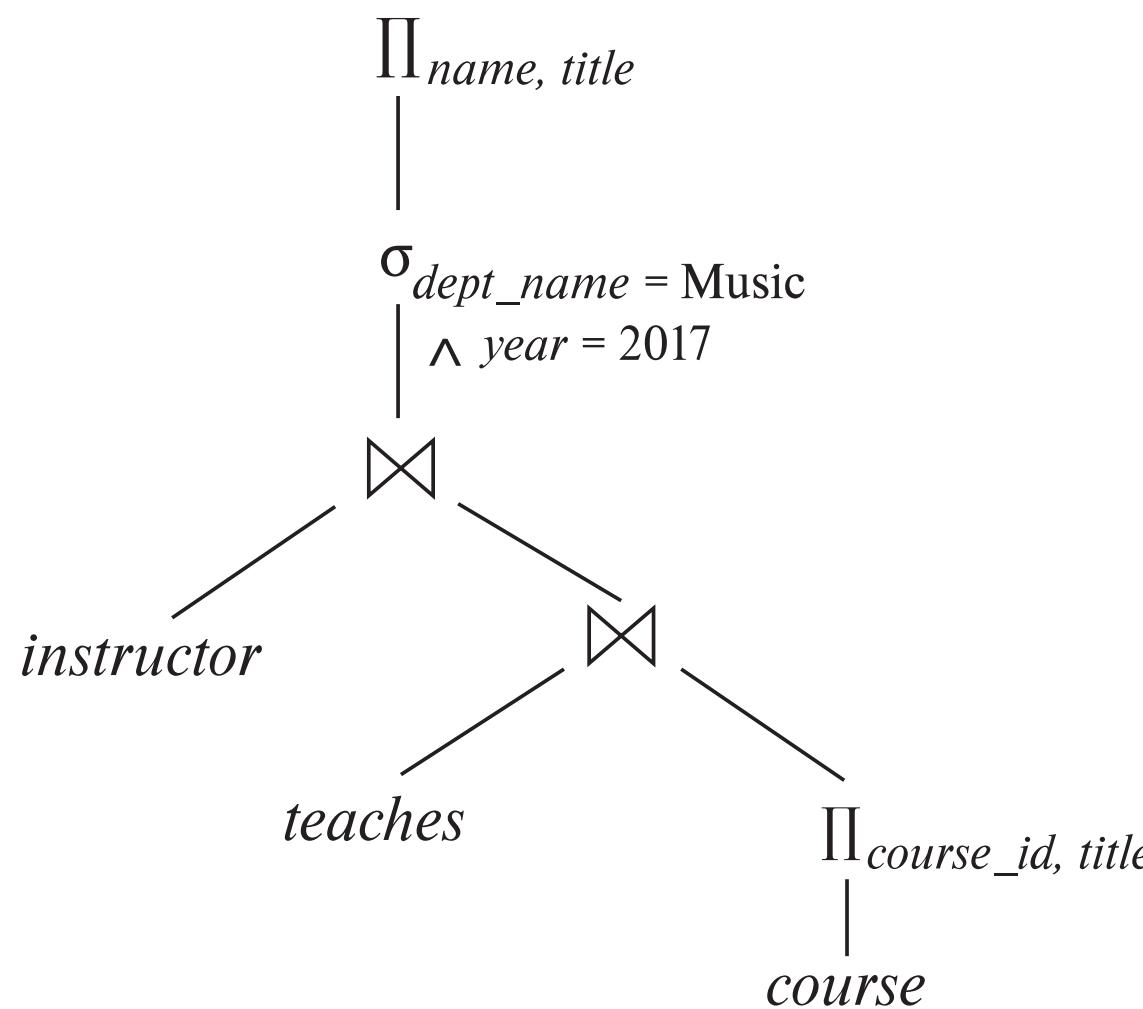
- $\Pi_{name, title}(\sigma_{dept_name= \text{Music}}(instructor \bowtie (teaches \bowtie \Pi_{course_id, title}(course))))$
- Transformation using rule 7(a):
 - $\Pi_{name, title}((\sigma_{dept_name= \text{Music}}(instructor)) \bowtie (teaches \bowtie \Pi_{course_id, title}(course)))$

Perform selection as early as possible reduces the size of the relation to be joined

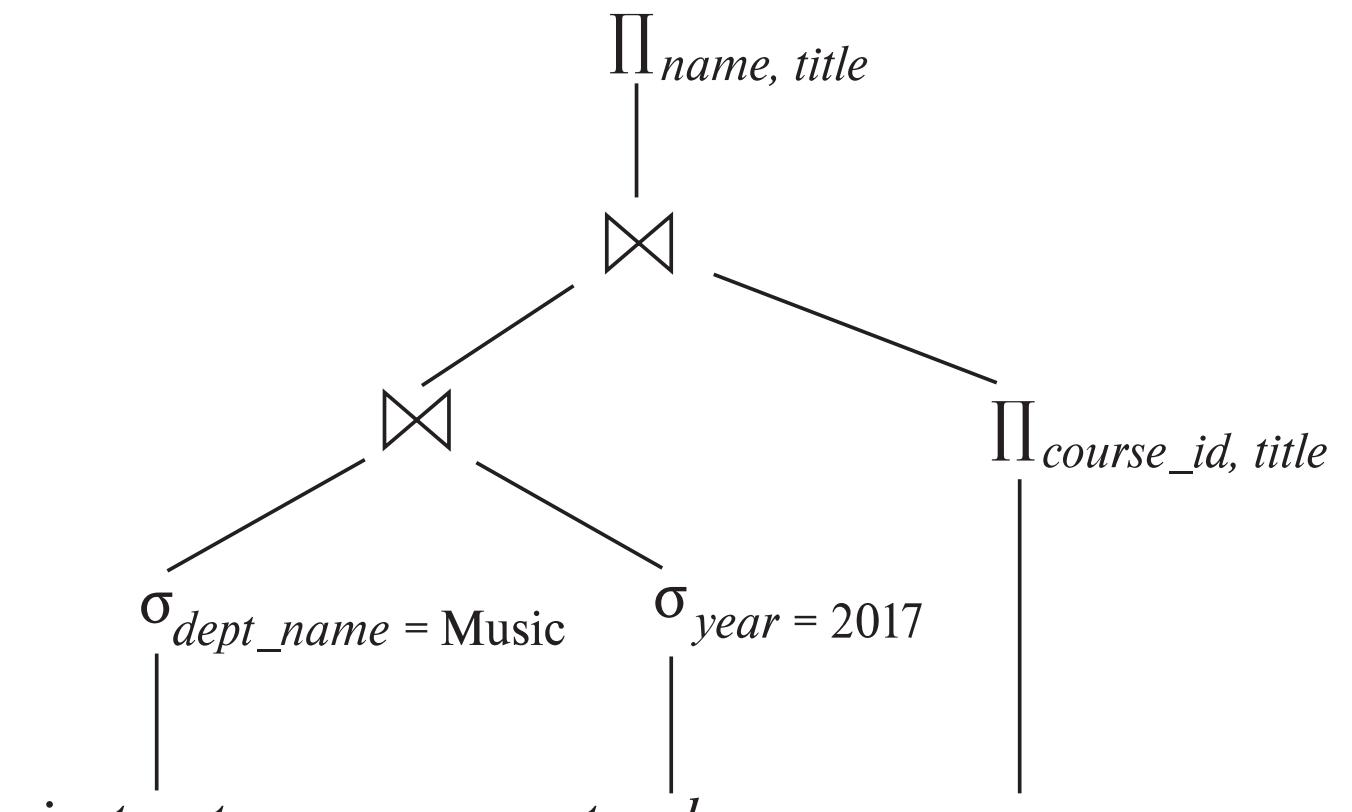
Transformation Example: Multiple Transformations

- Query: Find the names of all instructors in the Music department who have taught a course in 2017, along with the titles of the courses that they taught
 - $\Pi_{name, title}(\sigma_{dept_name = "Music" \wedge year = 2017} (instructor \bowtie (teaches \bowtie \Pi_{course_id, title} (course))))$
- Transformation using join associatively (Rule 6(a)):
 - $\Pi_{name, title}(\sigma_{dept_name = "Music" \wedge year = 2017} ((instructor \bowtie teaches) \bowtie \Pi_{course_id, title} (course)))$
- Second form provides an opportunity to apply the “perform selections early” rule, resulting in the subexpression:
 - $\sigma_{dept_name = "Music"} (instructor) \bowtie \sigma_{year = 2017} (teaches)$

Transformation Example: Multiple Transformations



(a) Initial expression tree



(b) Tree after multiple transformations

* Transformation Example: Pushing Projections

- Consider $\Pi_{name, title}((\sigma_{dept_name = "Music"}(instructor) \bowtie teaches) \bowtie \Pi_{course_id, title}(course))$

- When we compute

$$(\sigma_{dept_name = "Music"}(instructor) \bowtie teaches),$$

we obtain a relation whose schema is:

$$(ID, name, dept_name, salary, course_id, sec_id, semester, year)$$

- Push projections using equivalence rules 8a and 8b; eliminate unneeded attributes from intermediate results to get:

$$\Pi_{name, title}(\Pi_{name, course_id}(\sigma_{dept_name = "Music"}(instructor) \bowtie teaches) \bowtie \Pi_{course_id, title}(course))$$

Perform projections as early as possible reduces the size of the relation to be joined

Join Ordering Example

- For all relations r_1, r_2 , and r_3 ,

$$(r_1 \bowtie r_2) \bowtie r_3 = r_1 \bowtie (r_2 \bowtie r_3)$$

- * (Join Associativity) \bowtie
- If $r_2 \bowtie r_3$ is quite large and $r_1 \bowtie r_2$ is small, we choose

$$(r_1 \bowtie r_2) \bowtie r_3$$

so that we compute and store a smaller temporary relation

Cost Estimation

- Cost difference between evaluation plans for a query can be enormous
 - E.g., seconds vs. days in some cases
- Steps in cost-based query optimization
 - 1. Generate logically equivalent expressions using equivalence rules
 - 2. Annotate resultant expressions to get alternative query plans
 - 3. Choose the cheapest plan based on estimated cost

Cost Estimation

- Estimation of plan cost based on:
 - Statistical information about relations, such as:
 - number of tuples, number of distinct values for an attribute
 - Statistics estimation for intermediate results
 - to compute cost of complex expressions
 - Cost formulae for algorithms, computed using statistics

For more, please refer to Section 16.3 “Estimating Statistics of Expression Results” in the reference textbook